



Concepts and Models of Parallel and Data-centric Programming

Distributed Shared Memory

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline


- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory**
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. PGAS Foundations
 - b. DASH Overview
 - c. Distributed Data Structured
 - d. DASH Algorithms
 - e. Tasking

PGAS Foundations



Parallel Architectures


- Distributed Memory
 - Each processing element (P) has its separate main memory (M)



- Data exchange is achieved through message passing over the network

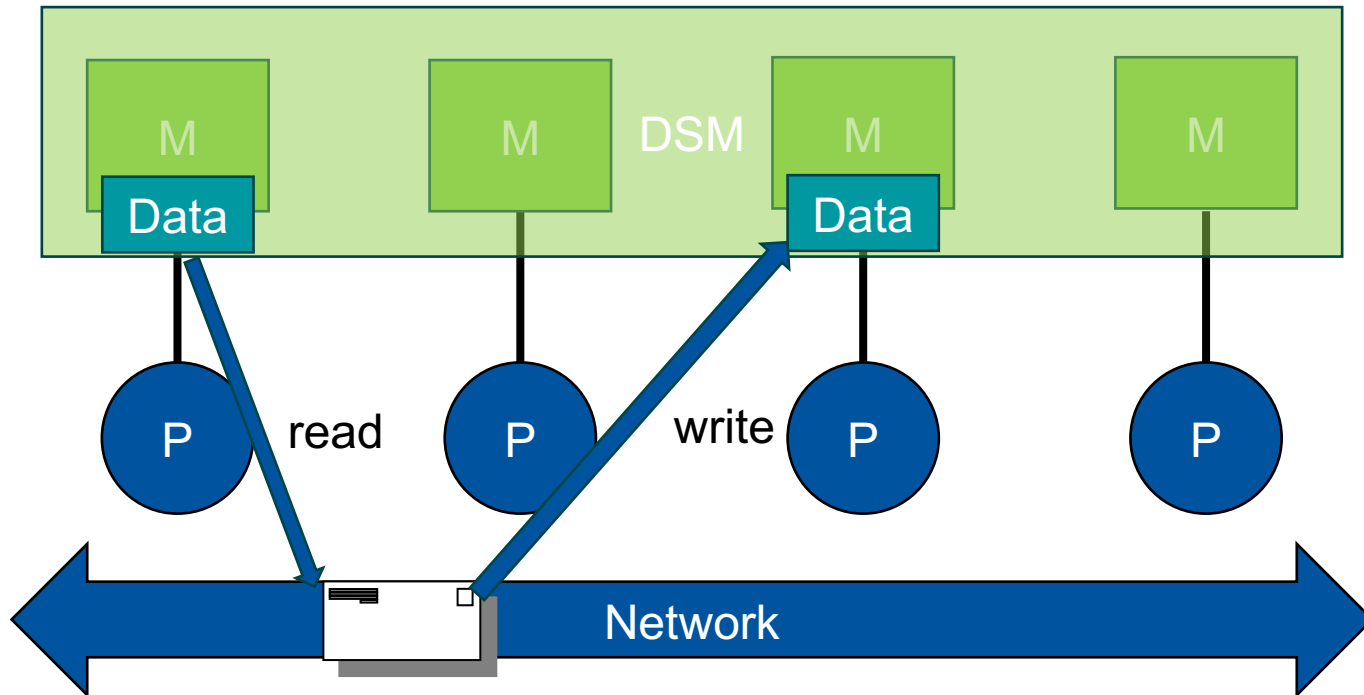
Parallel Architectures

- Distributed Shared Memory
 - All processing elements (P) have direct access to their main memory (M)
 - Each processing element (P) has access to the shared memory (DSM)




Parallel Architectures

- Distributed Shared Memory
 - All processing elements (P) have direct access to their main memory (M)
 - Each processing element (P) has access to the shared memory (DSM)




STREAM Triad

- Given: m -element vectors A , B , C
- Compute: $\forall i \in 1 \dots m: A_i = B_i + \alpha * C_i$
- In pictures:




STREAM Triad

- Given: m -element vectors A , B , C
- Compute: $\forall_{i \in 1 \dots m}: A_i = B_i + \alpha * C_i$
- In pictures, in parallel (multicore):




STREAM Triad

- Given: m -element vectors A , B , C
- Compute: $\forall_{i \in 1 \dots m}: A_i = B_i + \alpha * C_i$
- In pictures, in parallel (distributed memory):



STREAM Triad

- Given: m -element vectors A , B , C
- Compute: $\forall_{i \in 1 \dots m}: A_i = B_i + \alpha * C_i$
- In pictures, in parallel (distributed memory multicore):



HPC STREAM Triad: MPI + OpenMP

```
#include <hpcc.h>
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
        0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
        sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).
\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;

#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
}
```

Parallel Programming Models

- Shared Memory, e.g., Threading
 - Support dynamic, fine-grain parallelism
 - Considered simpler, more like traditional programming
 - If you want to access something, simply name it
- But:
 - Limited scalability
 - Bugs can be subtle, difficult to track down

Parallel Programming Models

- Distributed Memory, e.g., Message-Passing with MPI
 - More constrained model, can only access local data
 - Runs on all large-scale parallel platforms
 - And often can achieve near-optimal performance
 - Can serve as a strong foundation for high-level models
- But:
 - Communication must be used to get copies of remote data
 - Reveals (too) much about how to transfer, not what to transfer
 - Couples data transfer and synchronization
 - Has classes of bugs of its own

Parallel Programming Models

- Hybrid, e.g., MPI + Threading
 - Popular in the context of HPC (Threading: OpenMP) for highest performance
 - Division of labor: each handles what it does best
 - Overheads are amortized across processor cores
- But:
 - Requires multiple notations to express single logical parallel algorithm
 - Distinct semantics of MPI and Threading / OpenMP
 - May hold surprises in terms of unexpected side effects...

Parallel Programming Models

- Partitioned Global Address Space
 - (or: partitioned global namespace)
 - Can come as a language (extension) or API
 - Support a shared namespace on distributed memory ...
 - Permit any parallel task to access any lexically visible variable
 - ... with a sense of ownership
 - Variables have well-defined location, local accesses are cheaper than remote ones



Comparison

Name	Memory Model	Programming Model	Execution Model	Data Structures	Communication
MPI	Distributed Memory	SPMD	Cooperating Executables	Manually fragmented	APIs
OpenMP	Shared Memory	Global-view parallelism	Threading	Shared memory arrays	-
BSP	Distributed Memory	SPMD	SPMD	Distributed variables, Co-arrays	Explicit and Implicit
PGAS	Partitioned Global Address Space	SPMD	SPMD	Co-arrays	Implicit

- In this lecture, we present the DASH model
 - Asynchronous Partitioned Global Address Space
 - Implementation developed in the context of DFG SPPEXA prog