



Concepts and Models of Parallel and Data-centric Programming

Shared Memory II

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory**
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. Processes and Threads
 - b. Threading in C++
 - c. RAII idiom, Move Semantics
 - d. Mutual Exclusion
 - e. Condition Variable
 - f. Example: Queue

Threading in C++: Basics

Concurrency vs. Parallelism

- To many, these terms mean the same
 - Difference in focus and intent
- Both terms: running multiple tasks simultaneously
- Parallelism: performance-oriented
 - Taking advantage of hardware to increase performance
- Concurrency: separation of concern
 - Responsiveness

Threads – a programmer's view

- Execution stream within a program
 - Multiple threads working together may deliver a speedup
 - Programmer is responsible to distribute work
 - Programmer is responsible to manage the threads
- You can tell a thread:
 - What to do
 - When to start
- You can:
 - Wait for it to finish
- Other stuff (not relevant here):
 - Interrupt it, give it priority, etc.

Threads in C++ / 1

- Every C++ program has at least one thread running `main()`
 - Started by the C++ runtime
 - Additional threads run concurrently with each other, and the initial one
- Class `std::thread`
 - Representation of a “system thread”
 - Each thread has a function which is executed when it starts
 - Thread vanishes when the function returns
 - Defined in header `<thread>`
 - Reference: <https://en.cppreference.com/w/cpp/thread/thread>

Threads in C++ / 2

- After a thread has been started, the main thread could
 - wait for it to finish
 - `std::thread::join()`
 - blocks until the thread has finished execution
 - detach it
 - `std::thread::detach()`
 - permits the thread to execute independently from the thread handle
- Has to be handled before `std::thread` instance is destroyed
 - otherwise the `std::thread` destructor calls `std::terminate()`
- Data access by a thread has to be valid until it finishes

Threads in C++ / 3

- Thread identifiers are of type `std::thread::id`
- Obtained via
 - `get_id()` member function of a `std::thread` object
 - Current thread: `std::this_thread::get_id()`
- The id can be used to direct the control flow / divide work

```
1  std::thread::id master_thread;  
2  /* ... code ... */  
3  
4  if (std::this_thread::get_id() == master_thread)  
5      do_master_work();  
6  do_common_work();
```


Examples

Starting a thread / 1

- Starting a simple function as a thread

```
void do_some_work();  
std::thread my_thread(do_some_work);
```

- Starting a lambda expression as a thread

```
std::thread my_thread([]{  
    do_something();  
    do_something_else();  
});
```

Review: Lambdas in C++

- A lambda expression constructs a closure
 - unnamed function object capable of capturing variables in scope
 - Reference: <https://en.cppreference.com/w/cpp/language/lambda>
- Most commonly used syntax:
 - `[captures] (params) { body }`
 - `[captures] { body }`
 - return type can be derived from the return statement, or void if none
- Capture defaults:
 - `&`: capture by reference
 - `=`: capture by copy

Starting a thread / 2

- Using the ()-operator of a struct or class

```
1  class background_task
2  {
3  public:
4      void operator() () const
5      {
6          do_something();
7          do_something_else();
8      }
9  };
10
11  background_task f;
12  std::thread my_thread(f);
```

Starting a thread / 2

- Using the ()-operator of a struct or class

```
1  class background_task
2  {
3  public:
4      void operator() () const
5      {
6          do_something();
7          do_something_else();
8      }
9  };
10
11  background_task f;
12  std::thread my_thread(f);
```

Function call
operator

Possible error: undefined behavior

- Access to local (stack) data of the initial thread

```
1  struct func
2  {
3      int& i;
4      func(int& i_) : i(i_) {}
5      void operator() ()
6      {
7          for (unsigned j = 0; j < 1000000; ++j)
8              do_something(i);
9      }
10 };
11
12 void oops()
13 {
14     int some_local_var = 0;
15     func my_func(some_local_var);
16     std::thread my_thread(my_func);
17     my_thread.detach();
18 }
```

Possible error: undefined behavior

- Access to local (stack) data of the initial thread

```
1  struct func
2  {
3      int& i;
4      func(int& i_) : i(i_) {}
5      void operator() ()
6      {
7          for (unsigned j = 0; j < 1000000; ++j)
8              do_something(i);
9      }
10 };
11
12 void oops()
13 {
14     int some_local_var = 0;
15     func my_func(some_local_var);
16     std::thread my_thread(my_func);
17     my_thread.detach();
18 }
```

Thread might
still be
running

Possible error: undefined behavior

- Access to local (stack) data of the initial thread

```
1  struct func
2  {
3      int& i;
4      func(int& i_) : i(i_) {}
5      void operator() ()
6      {
7          for (unsigned j = 0; j < 1000000; ++j)
8              do_something(i);
9      }
10 };
11
12 void oops()
13 {
14     int some_local_var = 0;
15     func my_func(some_local_var);
16     std::thread my_thread(my_func);
17     my_thread.detach();
18 }
```

Potential
access to
dangling ref.

Thread might
still be
running

RAII idiom, Move Semantics

RAII idiom

- RAII: Resource Acquisition Is Initialization
 - binds the life cycle of a resource to the lifetime of an object
 - resource availability is a class invariant
 - guarantees the release of the resource in correct order
- Implementation of RAII:
 - encapsulate each resource into a class
 - constructor acquires the resource and establishes all class invariants or throws an exception if that cannot be done,
 - destructor releases the resource and never throws exceptions;
 - always use the resource via an instance of a RAII-class
 - with automatic storage duration or temporary lifetime

Using RAI to wait for completion

```
1  class thread_guard
2  {
3      std::thread& t;
4  public:
5      explicit thread_guard(std::thread& t_): t(t_) {}
6      ~thread_guard()
7      {
8          if (t.joinable())
9              t.join();
10     }
11     thread_guard(thread_guard const&) = delete;
12     thread_guard& operator= (thread_guard const&) = delete;
13 };
14
15 void f()
16 {
17     int some_local_var = 0;
18     func my_func(some_local_var);
19     std::thread t(my_func);
20     thread_guard g(t);
21 }
```

Using RAI to wait for completion

```
1  class thread_guard
2  {
3      std::thread& t;
4  public:
5      explicit thread_guard(std::thread& t_): t(t_) {}
6      ~thread_guard()
7      {
8          if (t.joinable())
9              t.join();
10     }
11     thread_guard(thread_guard const&) = delete;
12     thread_guard& operator= (thread_guard const&) = delete;
13 };
14
15 void f()
16 {
17     int some_local_var = 0;
18     func my_func(some_local_var);
19     std::thread t(my_func);
20     thread_guard g(t);
21 }
```

See above

Using RAI to wait for completion

```
1  class thread_guard
2  {
3      std::thread& t;
4  public:
5      explicit thread_guard(std::thread& t_): t(t_) {}
6      ~thread_guard()
7      {
8          if (t.joinable())
9              t.join();
10     }
11     thread_guard(thread_guard const&) = delete;
12     thread_guard& operator= (thread_guard const&) = delete;
13 };
14
15 void f()
16 {
17     int some_local_var = 0;
18     func my_func(some_local_var);
19     std::thread t(my_func);
20     thread_guard g(t);
21 }
```

See above

Destructor
will join if
possible and
necessary

Using RAI to wait for completion

```
1  class thread_guard
2  {
3      std::thread& t;
4  public:
5      explicit thread_guard(std::thread& t_): t(t_) {}
6      ~thread_guard()
7      {
8          if (t.joinable())
9              t.join();
10     }
11     thread_guard(thread_guard const&) = delete;
12     thread_guard& operator= (thread_guard const&) = delete;
13 };
14
15 void f()
16 {
17     int some_local_var = 0;
18     func my_func(some_local_var);
19     std::thread t(my_func);
20     thread_guard g(t);
21 }
```

See above

Destructor
will join if
possible and
necessary

g gets out of scope,
destructor is called

C++ move semantics

- *rvalue*, *lvalue*, and `&&`
 - An *lvalue* is an expression whose address can be taken. Anything you can make assignments to is an *lvalue*
 - An *rvalue* is an unnamed value that exists only during the evaluation of an expression.
 - The `&&` operator is like the reference operator (`&`), but whereas the `&` operator can only be used on *lvalues*, the `&&` operator can only be used on *rvalues*.
- It is possible to do a move (rather than a copy) if:
 - the object is an *rvalue*
 - the object's class defines the *special member move functions*
 - move constructor and move assignment operator

Ownership of threads

- `std::thread` employs the move semantic
 - moveable, but not copyable
 - because it is resource-owning
- Example:

```
1  void func1();
2  void func2();
3  std::thread t1(f1);
4  std::thread t2 = std::move(t1);
5  t1 = std::thread(f2);
6  std::thread t3;
7  t3 = std::move(t2);
8  t1 = std::move(t3);
```


Ownership of threads

- `std::thread` employs the move semantic
 - moveable, but not copyable
 - because it is resource-owning
- Example:

```
1  void func1();
2  void func2();
3  std::thread t1(f1);
4  std::thread t2 = std::move(t1);
5  t1 = std::thread(f2);
6  std::thread t3;
7  t3 = std::move(t2);
8  t1 = std::move(t3);
```

Program will
be terminated,
as t1 owns a
thread