



Concepts and Models of Parallel and Data-centric Programming

Shared Memory IV

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory**
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. Processes and Threads
 - b. Threading in C++
 - c. RAII idiom, Move Semantics
 - d. Mutual Exclusion
 - e. Condition Variable
 - f. Example: Queue

Example: Queue (part II: Challenges)

Queue: ADT design

- Container holding elements of type T
 - Function `enq()`: enqueue an element
 - Function `deq()`: de-queue an element
 - Internal representation of data could be an array
 - Variables `head` and `tail` point to indexes in that array

```
1  template<class T> class NaiveQueue
2  {
3  protected:
4      T* items; int head, tail;
5  public:
6      NaiveQueue() { ... initialization ... }
7      void enq(T element) { ... enqueue element: tail++ % QSIZE ... }
8      T deq() { ... return element: head++ % QSIZE ... }
9  };
```

Anything else?

- What may go wrong?
 - What if
 - Enqueuer finds a full array?
 - Dequeuer finds an empty array?
 - Wait for something to happen

Condition Variable

Motivation

- A thread could “actively” wait for an event to happen

```
1  bool flag;
2  std::mutex m;
3  void wait()
4  {
5      std::unique_lock<std::mutex> lk(m);
6      while (!flag)
7      {
8          lk.unlock();
9          std::this_thread::sleep_for(std::chrono::milliseconds(100));
10         lk.lock();
11     }
12 }
```

- What is “tricky” here?

Motivation

- A thread could “actively” wait for an event to happen

```
1  bool flag;
2  std::mutex m;
3  void wait()
4  {
5      std::unique_lock<std::mutex> lk(m);
6      while (!flag)
7      {
8          lk.unlock();
9          std::this_thread::sleep_for(std::chrono::milliseconds(100));
10         lk.lock();
11     }
12 }
```

Unlock

Lock (again)

- What is “tricky” here?

Condition / 1

- A condition variable implements notification
 - Block one (or multiple) thread(s) until another thread modified the condition variable and performs the notification
- Waiting thread(s):
 - Acquire the mutex (must be `std::unique_lock`)
 - Execute `wait()`, `wait_for()` or `wait_until()` (releases the mutex)
 - Wakeup occurs on the condition or after timeout, mutex is re-acquired
- Notifying thread:
 - Acquire the mutex () (usually via `std::lock_guard`)
 - Perform the modification, release the mutex
 - Execute `notify_one()` or `notify_all()`

Condition / 2

- Class `std::condition_variable`
 - Representation of a condition variable for “system threads”
 - Defined in header `<condition_variable>`
 - Reference: https://en.cppreference.com/w/cpp/thread/condition_variable
- Fundamentally, a condition variable is just an optimization over a busy-wait
 - Any code should work with a sub-optimal implementation
 - Repeated `unlock()` – `lock()` sequence

Example: Queue (part III: Implementation)

Thread-safe Queue / 1

- Thread-safe queue make use of `std::queue` internally
- Allows two (or more) threads to call `push()` and `wait_and_pop()` simultaneously
- Mutex, Condition variable, and implementation details are all hidden from the „user“

```
1  template<typename T> class threadsafe_queue
2  {
3  private:
4      std::queue<T> data;
5      std::mutex mut;
6      std::condition_variable cond;
7  public:
8      threadsafe_queue() {}
9
10 // continued on next slide
```

Thread-safe Queue / 2

```
1  public:
2      threadsafe_queue(threadsafe_queue const& other)
3      {
4          std::lock_guard<std::mutex> lk(other.mut);
5          data = other.data;
6      }
7      void push(T new_val)
8      {
9          std::lock_guard<std::mutex> lk(mut);
10         data.push(new_value);
11         cond.notify_one();
12     }
13     void wait_and_pop(T& value)
14     {
15         std::unique_lock<std::mutex> lk(mut);
16         cond.wait(lk, [this]{return ! data.empty() ;} );
17         value = data.front();
18         data.pop();
19     }
```

What you have learnt



Correct wording & Definitions / 1

- Critical Region:
 - A **section** of code that may only be executed by one process at any one time.
- Mutual Exclusion:
 - A **property**, the requirement that one thread of execution never enter its critical section at the same time that another concurrent thread of execution enters its own critical section.
- Starvation:
 - The **problem** encountered in concurrent computing where a process / thread is perpetually denied necessary resources to process its work.

Correct wording & Definitions / 2

- Deadlock:
 - A **program state** in which each member of a group (here: each thread) is waiting for some other member to take action.
- Race Condition:
 - Any situation where the outcome depends on the relative ordering of execution of operations on two or more threads
- Data Race: happens when there are two memory accesses in a program where both:
 - target the same location
 - are performed concurrently by two threads
 - are not reads
 - are not synchronization operations