



Concepts and Models of Parallel and Data-centric Programming

Shared Memory V

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory**
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- g. Futures
 - h. Example: QuickSort
 - i. Implementation of a Lock
 - j. Memory Consistency & Atomicity
 - k. Five Patterns of Synchronization

Futures



Futures / 1

- A future implements an (asynchronous) one-off event
 - Waiting for the occurrence of the event is possible
 - The event is typically the completion of an asynchronous task with return value
- Class `std::future` and `std::shared_future`
 - `std::future`: unique reference to the event
 - Move semantics
 - `std::shared_future`: multiple references from different threads to the same event possible (including multiple threads waiting)
 - Copy semantics
 - Calling `get()` will block until the (return) value has been computed
 - Defined in header `<future>`
 - Reference: <https://en.cppreference.com/w/cpp/thread/future>

Futures / 2

- Class `std::promise`
 - Defined in header `<future>`
 - Implementation detail: promise to set the value in the future
 - Reference: <https://en.cppreference.com/w/cpp/thread/promise>
- Function `std::async`
 - Runs a function asynchronously, returns a `std::future`
 - Reference: <https://en.cppreference.com/w/cpp/thread/async>

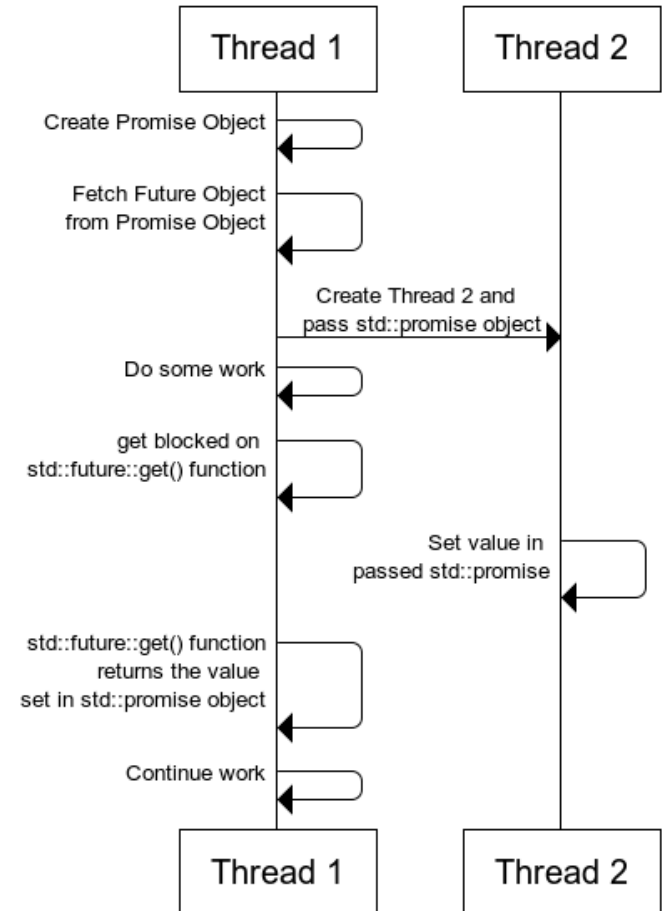
Futures / 3

- Illustration of an asynchronous execution with `std::future`

```
// let T be some valid typename  
std::promise<T> promiseObj;  
std::future<T> futureObj =  
    promiseObj.get_future();  
std::thread th( f , &promiseObj);  
th.join();
```

- Remark: in many cases, the `std::promise` object is not used explicitly

std::promise and std::future work flow

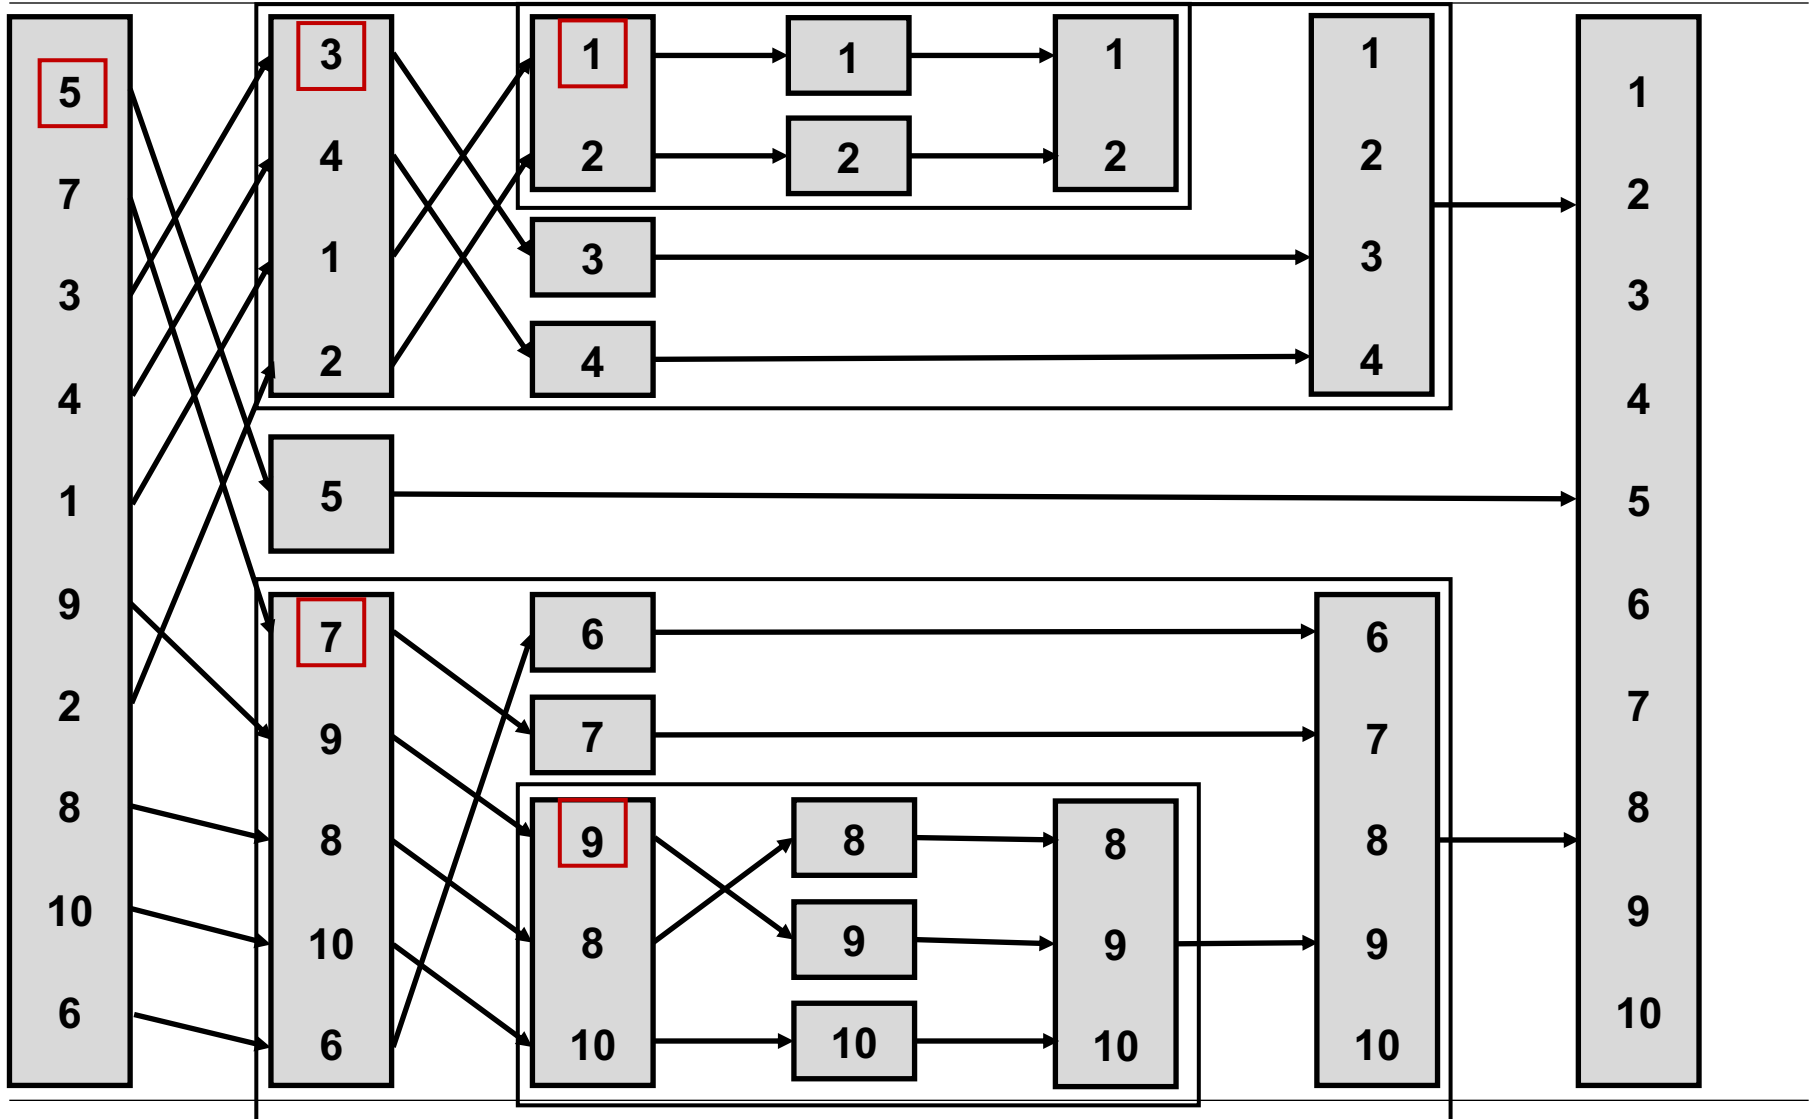


Example: QuickSort

Motivation

- FP-style parallel programming
 - Functional Programming (FP) in this context: the result of a function call depends solely on the parameters to that function
 - does not depend on any external state
 - trivial example: math routines from the standard library, such as `sin()`
- A *pure* function does not modify any external state; its effects are entirely limited to the return value
- Concurrency in functional programming may eliminate data races

FP-style QuickSort / 1



FP-style QuickSort / 2

```
1  template<typename T>
2  std::list<T> seq_qsort(std::list<T> input)
3  {
4      if(input.empty())
5          return input;
6
7      std::list<T> result;
8      result.splice(result.begin(), input, input.begin());
9      T const& pivot = *result.begin();
10
11     auto divide_point = std::partition(input.begin(), input.end(),
12         [&](T const& t){return t < pivot;});
13     std::list<T> lower_part;
14     lower_part.splice(lower_part.end(), input, input.begin(), divide_point);
15     auto new_lower(seq_qsort(std::move(lower_part)));
16     auto new_higher(seq_qsort(std::move(input)));
17     result.splice(result.end(), new_higher);
18     result.splice(result.begin(), new_lower);
19     return result;
20 }
```

FP-style QuickSort / 3

```
1  template<typename T>
2  std::list<T> seq_qsort(std::list<T> input)
3  {
4      if(input.empty())
5          return input;
6
7      std::list<T> result;
8      result.splice(result.begin(), input, input.begin());
9      T const& pivot = *result.begin();
10
11     auto divide_point = std::partition(input.begin(), input.end(),
12         [&](T const& t){return t < pivot;});
13     std::list<T> lower_part;
14     lower_part.splice(lower_part.end(), input, input.begin(), divide_point);
15     auto new_lower(seq_qsort(std::move(lower_part)));
16
17     auto new_higher(seq_qsort(std::move(input)));
18     result.splice(result.end(), new_higher);
19     result.splice(result.begin(), new_lower);
20     return result;
21 }
```

FP-style QuickSort / 4

```
1  template<typename T>
2  std::list<T> par_qsort(std::list<T> input)
3  {
4      if(input.empty())
5          return input;
6
7      std::list<T> result;
8      result.splice(result.begin(), input, input.begin());
9      T const& pivot = *result.begin();
10
11     auto divide_point = std::partition(input.begin(), input.end(),
12         [&](T const& t){return t < pivot;});
13     std::list<T> lower_part;
14     lower_part.splice(lower_part.end(), input, input.begin(), divide_point);
15     std::future<std::list<T> > new_lower(
16         std::async(&par_qsort<T>, std::move(lower_part)));
17     auto new_higher(par_qsort(std::move(input)));
18     result.splice(result.end(), new_higher);
19     result.splice(result.begin(), new_lower);
20     return result;
21 }
```

FP-style QuickSort / 5

- Remarks on the simple sequential implementation:
 - returns a list by-value (`std::sort` performs an in place sort)
 - interface is FP-style, implementation is partly imperative for efficiency reasons
 - `std::list`: first element is taken as pivot element
 - may result in sub-optimal sort
 - reason: avoid list traversal
- Bad parallelization: `std::partition` does a lot of work and is still sequential
- Unclear scalability: execution depends on quality of the implementation
 - how many threads will be created?