



Concepts and Models of Parallel and Data-centric Programming

Shared Memory VI

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory**
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- g. Futures
 - h. Example: QuickSort
 - i. Implementation of a Lock
 - j. Memory Consistency & Atomicity
 - k. Five Patterns of Synchronization

Implementation of a lock

First approach: naïve lock

```
1 bool flag[2];    // initialization w/ false
2 void lock() {
3     flag[i] = true;
4     while (flag[j]) {}
5 }
```

- Illustration for two threads (*i* and *j*) only
 - Each thread has a flag
 - *i* = 0, *j* = 1
 - `unlock()`: set thread's flag back to false
- Would this work?

First approach: naïve lock

```
1 bool flag[2];    // initialization w/ false
2 void lock() {
3     flag[i] = true;
4     while (flag[j]) {}
5 }
```

set flag

Wait for other flag to become false

- Illustration for two threads (i and j) only
 - Each thread has a flag
 - $i = 0, j = 1$
 - `unlock()`: set thread's flag back to false
- Would this work?

Peterson's Algorithm

- Lock algorithm
 - that does not fail when concurrent
 - that does not deadlock when sequential

```
1  void lock() {
2    flag[i] = true;
3    victim = i;
4    while (flag[j] && victim == i) {};
5  }
6  void unlock() {
7    flag[i] = false;
8  }
```

Peterson's Algorithm

- Lock algorithm
 - that does not fail when concurrent
 - that does not deadlock when sequential

```
1  void lock() {  
2  flag[i] = true;   
3  victim = i;  
4  while (flag[j] && victim == i) {};  
5  }  
6  void unlock() {  
7  flag[i] = false;  
8  }
```

Announce I am interested

Peterson's Algorithm

- Lock algorithm
 - that does not fail when concurrent
 - that does not deadlock when sequential

```
1  void lock() {  
2  flag[i] = true;  
3  victim = i;  
4  while (flag[j] && victim == i) {};  
5  }  
6  void unlock() {  
7  flag[i] = false;  
8  }
```

Announce I am interested

Defer to other

Peterson's Algorithm

- Lock algorithm
 - that does not fail when concurrent
 - that does not deadlock when sequential

```
1  void lock() {  
2  flag[i] = true;  
3  victim = i;  
4  while (flag[j] && victim == i) {};  
5  }  
6  void unlock() {  
7  flag[i] = false;  
8  }
```

Announce I am interested

Defer to other

Wait while other is interested and I am the victim

Peterson's Algorithm

- Lock algorithm
 - that does not fail when concurrent
 - that does not deadlock when sequential

```
1  void lock() {  
2    flag[i] = true;  
3    victim = i;  
4    while (flag[j] && victim == i) {};  
5  }  
6  void unlock() {  
7    flag[i] = false;  
8  }
```

Announce I am interested

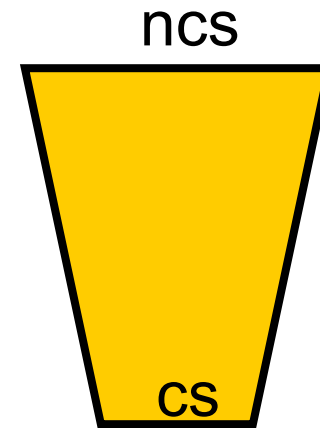
Defer to other

Wait while other is interested and I am the victim

I am no longer interested

Remark on generalization

- Generalization for Peterson's Algorithms for n threads: Filter Alg.
 - The Filter lock creates $n-1$ "waiting rooms" (levels) that a thread must traverse before acquiring the lock
 - At each level
 - At least one enters level
 - At least one blocked if many try
 - Only one thread makes it through
- Issue of Filter Alg.: Fairness
 - A thread can be overtaken arbitrary number of times
 - One solution: bakery algorithm with first come first serve property
 - But impractical efficiency!!!



Test-and-set locks

Overview

- The previous examination of a lock protocol was
 - helpful in understanding the foundations
 - meant to illustrate the underlying problems

Overview

- The previous examination of a lock protocol was
 - helpful in understanding the foundations
 - meant to illustrate the underlying problems
- However, in the “real” world, these protocols are not efficient
- Modern processors: **MIMD** (Multiprocessors)
 - Multiple instruction, multiple data
 - Provide a test-and-set instruction ...
 - ... that can be used to implements locks efficiently

Test-and-set Instruction

- Modern architectures provide a test-and-set instruction
 - Write 1 to a memory location
 - Return its old value
 - Both is done as a single *atomic* operation

Test-and-set Instruction

- Modern architectures provide a test-and-set instruction
 - Write 1 to a memory location
 - Return its old value
 - Both is done as a single *atomic* operation
- Atomic Operation:
 - A **set of operations** that appears to the rest of the system to occur at once without being interrupted (it is uninterruptable and indivisible).

Test-and-set Instruction

- Modern architectures provide a test-and-set instruction
 - Write 1 to a memory location
 - Return its old value
 - Both is done as a single *atomic* operation
- Atomic Operation:
 - A **set of operations** that appears to the rest of the system to occur at once without being interrupted (it is uninterruptable and indivisible).
- General: Test-and-set (TAS)
 - Swap **true** with current value
 - Return value tells if prior value was **true** or **false**

Test-and-set Lock

- Locking
 - Lock is free: value is false
 - Lock is taken: value is true
- Acquire lock by calling TAS
 - If result is false, you win
 - If result is true, you lose
- Release lock by writing false
- Why is this more efficient than the Peterson Algorithm?

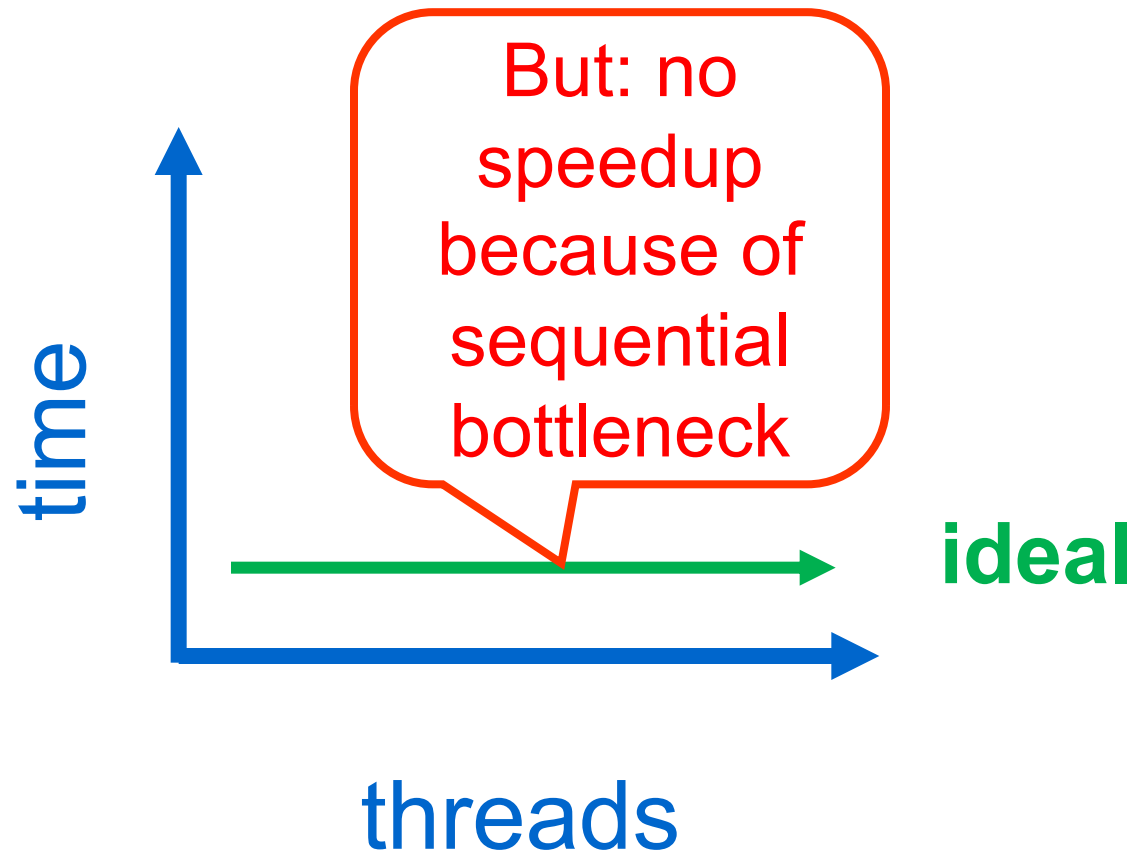
Test-and-set Lock

- Locking
 - Lock is free: value is false
 - Lock is taken: value is true
- Acquire lock by calling TAS
 - If result is false, you win
 - If result is true, you lose
- Release lock by writing false

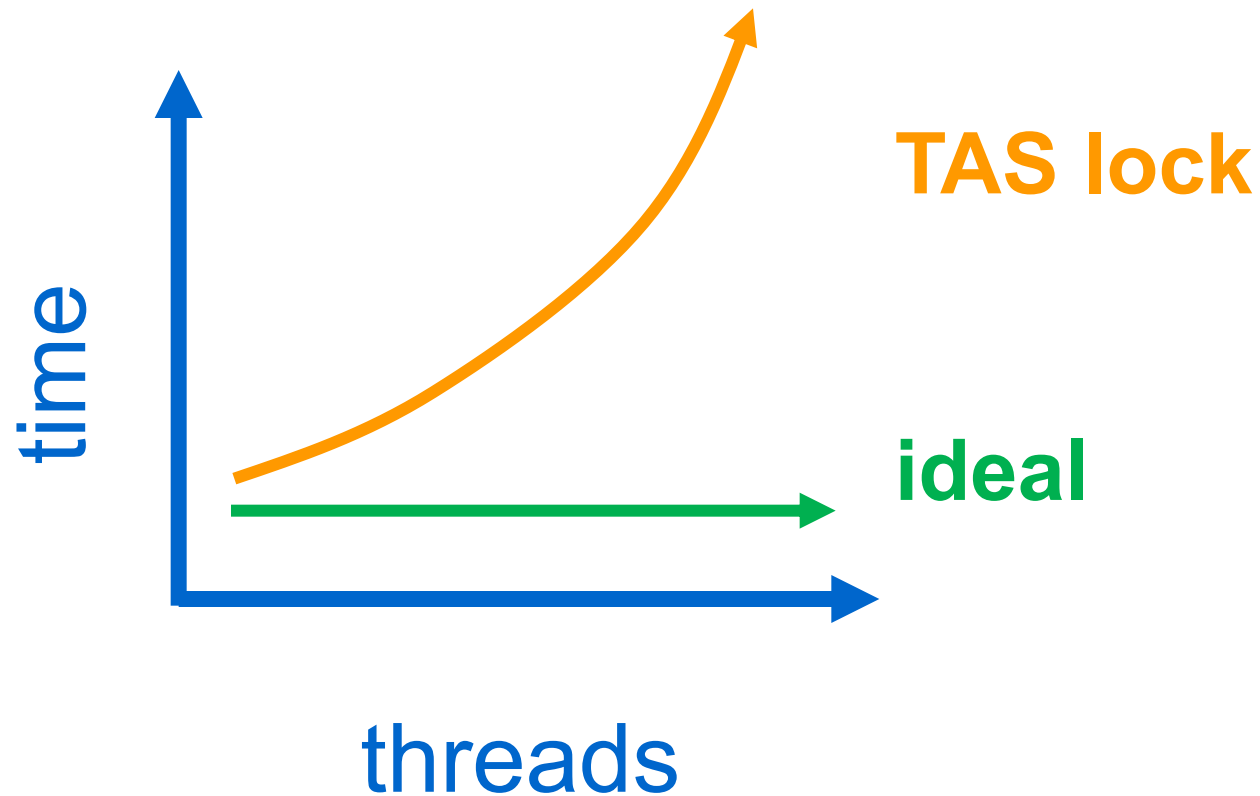
- Why is this more efficient than the Peterson Algorithm?
- Because: $O(1)$ space / instances (as opposed to $O(n)$)!

Remark: do-it-yourself vs. standard implementations

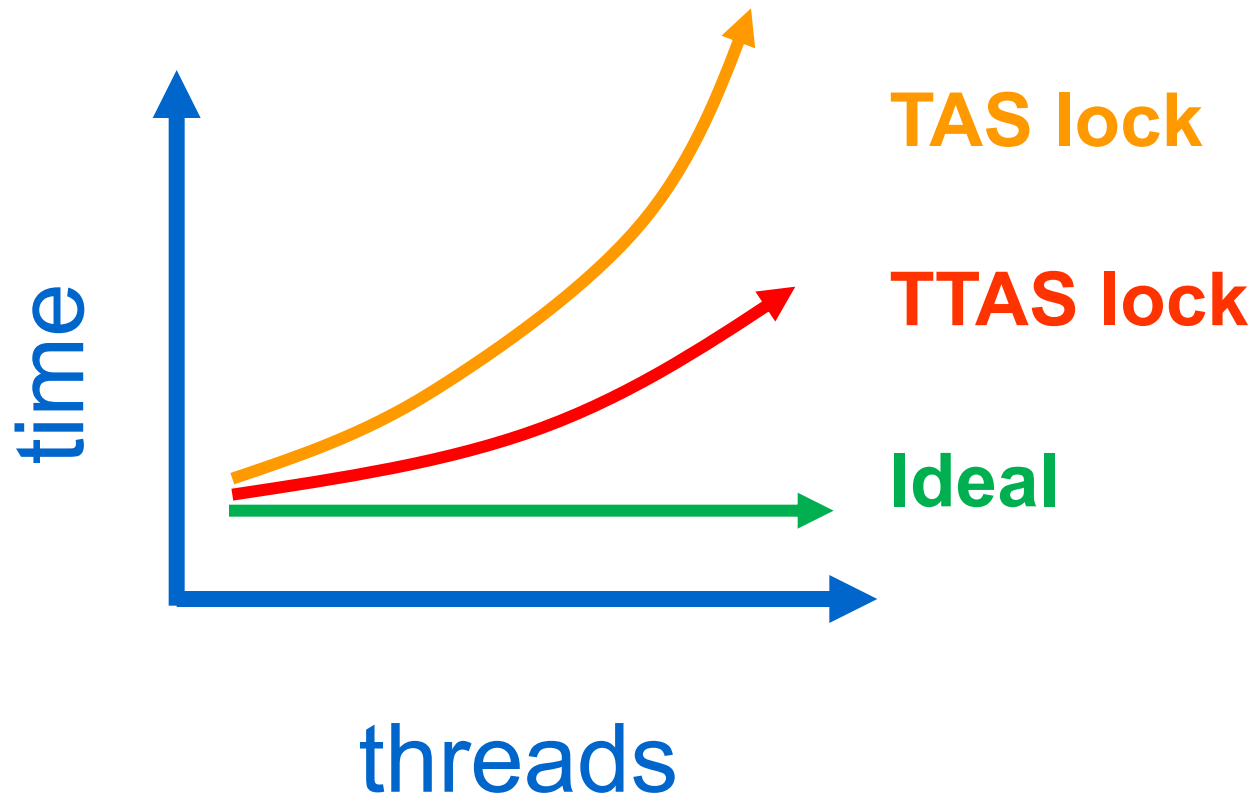
Performance Considerations / 1



Performance Considerations / 2



Performance Considerations / 3



- TTAS: Locking is split up into two phases:
 - Waiting until the lock is available ... before trying to acquire the lock

Simplified Explanation

- TAS invalidates cache lines
 - Spinners
 - Miss in cache
 - Go to bus
 - Thread wants to release lock: delayed behind spinners

Simplified Explanation

- TAS invalidates cache lines
 - Spinners
 - Miss in cache
 - Go to bus
 - Thread wants to release lock: delayed behind spinners
- TTAS waits until lock “looks” free
 - Spin on local cache
 - No bus use while lock busy
 - Problem: when lock is released: invalidation storm ...

Simplified Explanation

- TAS invalidates cache lines
 - Spinners
 - Miss in cache
 - Go to bus
 - Thread wants to release lock: delayed behind spinners
- TTAS waits until lock “looks” free
 - Spin on local cache
 - No bus use while lock busy
 - Problem: when lock is released: invalidation storm ...
- Solution 1: learn (much) more about system architecture

Simplified Explanation

- TAS invalidates cache lines
 - Spinners
 - Miss in cache
 - Go to bus
 - Thread wants to release lock: delayed behind spinners
- TTAS waits until lock “looks” free
 - Spin on local cache
 - No bus use while lock busy
 - Problem: when lock is released: invalidation storm ...
- Solution 1: learn (much) more about system architecture
- Solution 2: use lock methods provided by runtime system 😊