



# Concepts and Models of Parallel and Data-centric Programming

Shared Memory IX

Lecture, Summer 2020

Dr. Christian Terboven <[terboven@itc.rwth-aachen.de](mailto:terboven@itc.rwth-aachen.de)>

# Outline

---


- 0. Organization
  - 1. Foundations
  - 2. Shared Memory**
  - 3. GPU Programming
  - 4. Bulk-Synchronous Parallelism
  - 5. Message Passing
  - 6. Distributed Shared Memory
  - 7. Parallel Algorithms
  - 8. Parallel I/O
  - 9. MapReduce
  - 10. Apache Spark
- l. Coarse-grained Synchronization
  - m. Fine-grained Synchronization
  - n. Optimistic Synchronization
  - o. Lazy Synchronization
  - p. Lock-free Synchronization

# Linked List: coarse-grained synchronization

# Coarse-grained Synchronization


---

- Each method locks the object
  - Avoid contention with optimized lock type
  - Easy to reason about




# Illustration: add

- Simple ...




# Illustration: add

- Simple ...




- ... but: bottleneck!




# Illustration: add

- Simple ...



- ... but: bottleneck!



# Summary: Coarse-grained Sync.

---

- Sequential bottleneck
  - Threads “stand in line”
- Adding more threads
  - Does not improve throughput
  - Struggle to keep it from getting worse
- So why even use a multiprocessor?
  - Simple, clearly correct - deserves some respect!



# Linked List: fine-grained synchronization



# Fine-grained Synchronization

---


- Instead of using a single lock: split object into
  - Independently-synchronized components
  - Methods conflict when they access the same component at the same time
- Improve concurrency by locking individual entries
  - instead of placing a lock on the entire list, add a lock to each entry
  - as a thread traverses the list, it locks each entry when it first visits, and sometime later releases it
  - Methods that work on disjoint pieces need not exclude each other
- Requires careful thought

“Do not meddle in the affairs of wizards, for they are subtle and quick to anger”


# First try of concurrent remove / 1




## First try of concurrent remove / 2




## First try of concurrent remove / 3




## First try of concurrent remove / 4



# First try of concurrent remove / 5




# First try of concurrent remove / 6







# First try of concurrent remove / 7



## First try of concurrent remove / 8




# First try of concurrent remove / 9




## Problem of first try

- To delete node c
  - Swing node b's next field to d
- Problem is,
  - Someone deleting b concurrently could direct a pointer to c



## Problem of first try

- To delete node c
  - Swing node b's next field to d
- Problem is,
  - Someone deleting b concurrently could direct a pointer to c
- If a node is locked
  - No one can delete node's *successor*
- If a thread locks
  - Node to be deleted
  - And its predecessor
  - Then it works



# Concurrent remove / 1

- Boilerplate

```
1  bool Node::remove(int item_key)
2  {
3      Node *pred, *curr;
4      std::unique_lock<std::mutex> lpred(pred->mut,
5          std::defer_lock);
6      std::unique_lock<std::mutex> lcurr(curr->mut,
7          std::defer_lock);
8      {
9          ...
10     }
11 }
```

Ensure that locks  
are released

Everything else

- Note: a `std::mutex` has to be added to the Node class

# Concurrent remove / 1

- Boilerplate

```
1  bool Node::remove(int item_key)
2  {
3      Node *pred, *curr;
4      std::unique_lock<std::mutex> lpred(pred->mut,
5          std::defer_lock);
6      std::unique_lock<std::mutex> lcurr(curr->mut,
7          std::defer_lock);
8      {
9          ...
10     }
11 }
```

Ensure that locks  
are released

Everything else


- Note: a `std::mutex` has to be added to the Node class
- Q: why do we have to ensure that locks are released?

## Concurrent remove / 2

- Remove method

```
pred = this->head;  
lpred.lock();  
curr = pred->next;  
lcurr.lock();  
...
```

Lock pred == head





## Concurrent remove / 2

- Remove method


```
pred = this->head;  
lpred.lock();
```

```
curr = pred->next;  
lcurr.lock();
```

```
...
```

Lock current

List traversal




# Concurrent remove / 3

- List traversal

```
while (curr->key <= item_key) {  
    if (item_key == curr->key) {  
        pred->next = curr->next;  
        return true;  
    }  
    lpred.unlock();  
    pred = curr;  
    curr = curr->next;  
    lcurr.lock();  
}  
return false;
```

At start of each loop:  
curr and pred locked




# Concurrent remove / 3

- List traversal

```
while (curr->key <= item_key) {  
    if (item_key == curr->key) {  
        pred->next = curr->next;  
        return true;  
    }  
    lpred.unlock();  
    pred = curr;  
    curr = curr->next;  
    lcurr.lock();  
}  
return false;
```

If item found:  
remove




# Concurrent remove / 3

- List traversal

```
while (curr->key <= item_key) {  
    if (item_key == curr->key) {  
        pred->next = curr->next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr->next;  
    curr.lock();  
}  
return false;
```

Proceed



## Concurrent remove / 3

- List traversal

```
while (curr->key <= item_key) {  
    if (item_key == curr->key) {  
        pred->next = curr->next;  
        return true;  
    }  
    lpred.unlock();  
    pred = curr;  
    curr = curr->next;  
    lcurr.lock();  
}  
return false;
```

Element not present

# Adding an element

---

- To add node e
  - Must lock predecessor
  - ~~Must lock successor~~
- Neither can be deleted

# Summary: Fine-grained Sync.

---

- Easy to check that
  - tail always reachable from head
  - nodes sorted, no duplicates
- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Inefficient