



# Concepts and Models of Parallel and Data-centric Programming

Shared Memory XII

Lecture, Summer 2020

Dr. Christian Terboven <[terboven@itc.rwth-aachen.de](mailto:terboven@itc.rwth-aachen.de)>

# Outline

---

- 0. Organization
  - 1. Foundations
  - 2. Shared Memory**
  - 3. GPU Programming
  - 4. Bulk-Synchronous Parallelism
  - 5. Message Passing
  - 6. Distributed Shared Memory
  - 7. Parallel Algorithms
  - 8. Parallel I/O
  - 9. MapReduce
  - 10. Apache Spark
- o. Lock-free Synchronization
  - p. SIMD / Vectorization
  - q. Intrinsics for SIMD
  - r. Parallel STL for SIMD and Parallelism

---

# Review

# Set Interface

---

- Unordered collection of items
  - No duplicates
- Methods
  - `add(x)` put `x` in set
  - `remove(x)` take `x` out of set
  - `contains(x)` tests if `x` in set

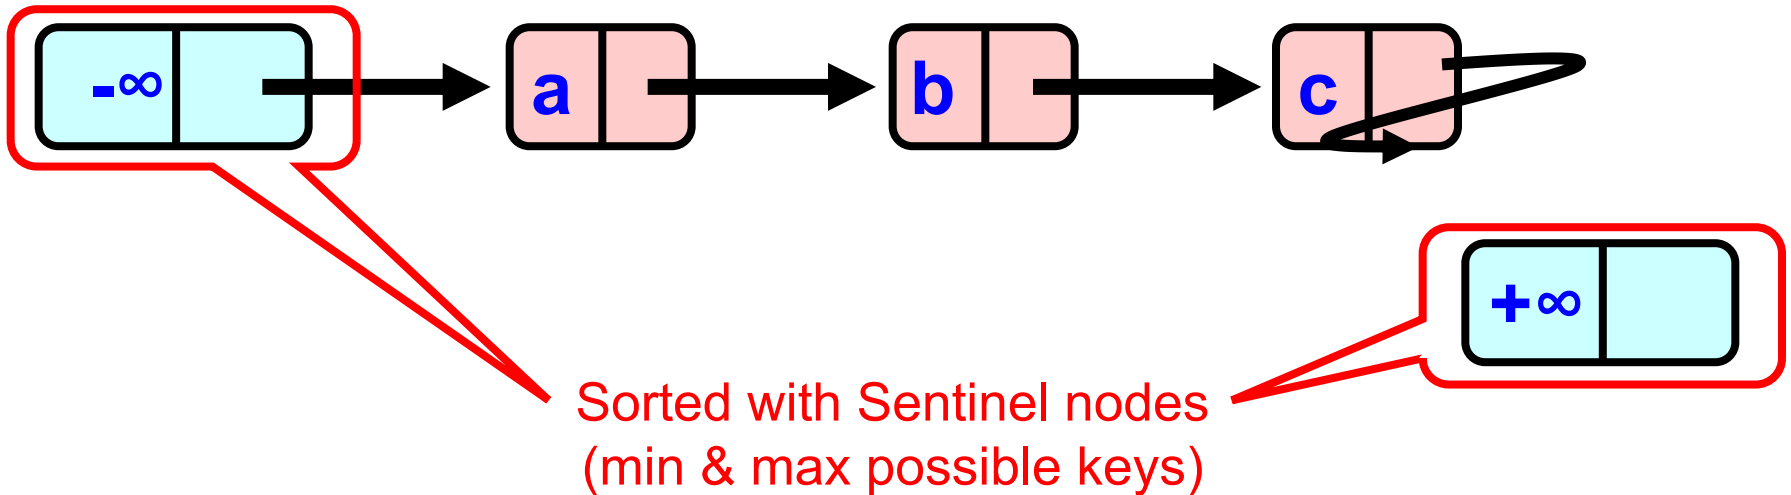
```
1  template<typename T> class Set {  
2  public:  
3      bool add(T x);  
4      bool remove(T x);  
5      bool contains(T x);  
6  };
```

# List-based Set

- List node:

```
1  template<typename T> class Node {  
2  public:  
3      T item;  
4      int key;  
5      Node* next;  
6  };
```

Item of interest  
Hash  
Reference to next node



# Comparison

---

	<b>add()</b>	<b>remove()</b>	<b>contains()</b>
Coarse-grained Sync.	whole object locked	whole object locked	whole object locked
Fine-grained Sync.	chain of pair-wise acquire and release	chain of pair-wise acquire and release	no lock
Optimistic Sync.	lock targets only, but validate with traversal	lock targets only, but validate with traversal	chain of acquire and release
Lazy Sync.	mark and lock targets only, retry if conflict	mark and lock targets only, retry if conflict	no lock (check for marker)
Lock-free Sync.	no lock	no lock	no lock

---

# Lock-free Synchronization

# Goal of lock-free data structure

---

- No matter what ...
  - Guarantees minimal progress in any execution
  - i.e. some thread will always complete a method call
  - Even if others halt at malicious times
- Implies that implementation can't use locks
- Lock-free List
  - Next logical step
    - Wait-free contains()
    - lock-free add() and remove()
  - Use only `atomic` functionality

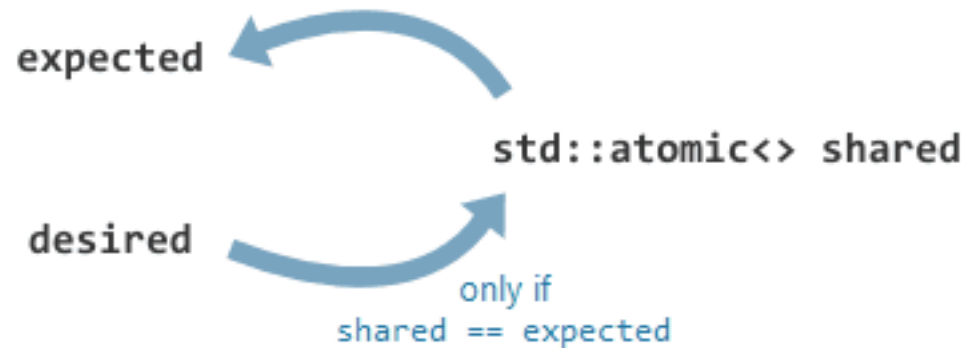




# Implementation sketch


---

- Use `std::atomic<>::compare_exchange_weak()`
  - Use of CAS: compare-and-swap atomic instruction
    - compares the value of a memory location with a given value and, only if they are the same, modifies the content of that memory location to a new given value
- Atomically
  - Swing reference and
  - Update marker
- Remove in two steps
  - Set mark bit in next field
  - Redirect predecessor's pointer



## Problem (if CAS used only for pointer)


---




## Problem (if CAS used only for pointer)

---


Logical Removal



## Problem (if CAS used only for pointer)



## Problem (if CAS used only for pointer)



Problem: Pointer modification and marker update are not considered together.

# Implementation sketch

---



Marker and Pointer  
are CASed together

# Implementation sketch

---


Logical Removal =  
Set Mark Bit



Marker and Pointer  
are CASed together

# Implementation sketch

Logical Removal =  
Set Mark Bit




Marker and Pointer  
are CASed together

Fail CAS: Node not  
added after logical  
Removal



# Implementation sketch

Logical Removal =  
Set Mark Bit




Marker and Pointer  
are CASed together


Physical  
Removal  
CAS

Fail CAS: Node not  
added after logical  
Removal


# Illustration: removal




# Illustration: removal




# Illustration: removal




# Illustration: removal



# Illustration: removal





# Illustration: removal



# Illustration: traversal


---



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)




# Illustration: traversal



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)


# Illustration: traversal

---




- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)

# Illustration: traversal



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)


# Illustration: traversal



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)

# Illustration: traversal


---



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)


# Illustration: traversal

---




- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)

# Illustration: traversal



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)


## Illustration: traversal



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)




## Illustration: traversal



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)

# Illustration: traversal


---



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)


# Illustration: traversal

---



- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)

# Illustration: traversal




- Reaction to a logically deleted node? CAS the predecessor's next field and proceed (repeat if necessary)

# Performance


---

- As presented in the book (Java, The Art of Multiprocessor Prog.):
  - Different list-based set implementations
  - 16-node machine
  - Vary percentage of `contains()` calls
- **Thesis offer:** Reproduce these measurements on current hardware
  - Inclusion of Transactional Memory
  - Sensitivity to the Memory Hierarchy

# Performance: high contains ratio



# Performance: low contains ratio



# Summary: Lock-free Sync.

---

- Don't use locks at all
  - Use `std::atomic<>::compare_exchange_weak()` ...
- Advantages
  - No Scheduler Assumptions/Support
- Disadvantages
  - Complex and more implementation work
  - Sometimes high overhead
- Locking vs. Non-blocking:
  - Extremist views on both sides
  - Remember: Blocking/non-blocking is a property of a method