



Concepts and Models of Parallel and Data-centric Programming

BSP I

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>


Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism**
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. Motivation
 - b. BSP Computer
 - c. BSP Programming Model
 - d. BSP Cost Model
 - e. Bulk Library

Motivation


Parallel Architectures

- HPC market is at large dominated by distributed memory *multicomputers*: *clusters* and specialised *supercomputers*
- Nodes have no direct access to other nodes' memory and run a separate copy of the (possibly stripped down) OS



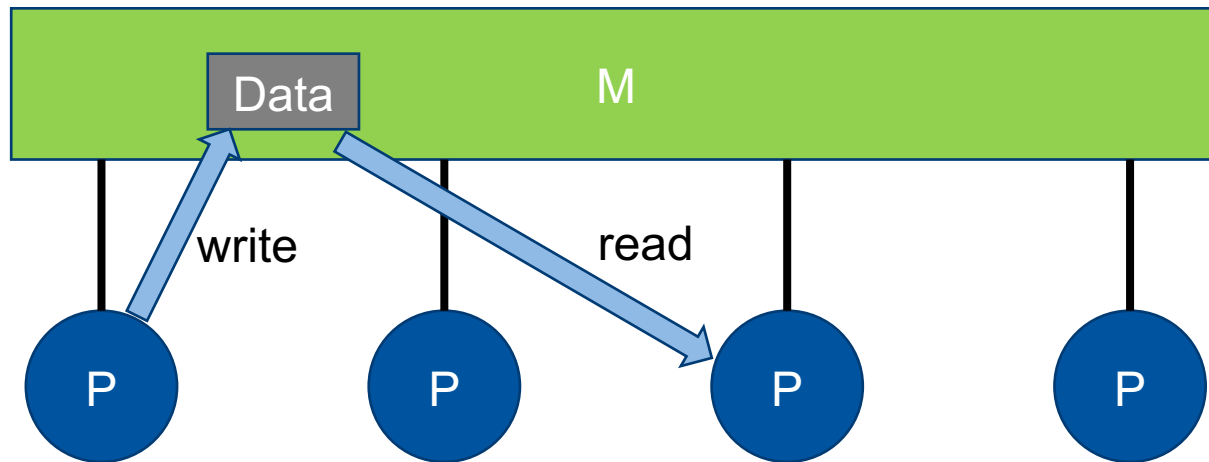
Parallel Architectures

- Shared Memory
 - All processing elements (P) have direct access to the main memory block (M)



Parallel Architectures


- Shared Memory
 - All processing elements (P) have direct access to the main memory block (M)



- Data exchange is achieved through read/write operations on shared variables located in the global address space


Parallel Architectures

- Distributed Memory
 - Each processing element (P) has its separate main memory block (M)



Parallel Architectures – Two-Sided Communication


- Distributed Memory
 - Each processing element (P) has its separate main memory block (M)



- Data exchange is achieved through **message** passing over the network
- **Both** processes actively call send and receive function

Parallel Architectures – One-Sided Communication

- Distributed Memory
 - Each processing element (P) has its separate main memory block (M)



- Alternative: Remote Direct Memory Access (RDMA)
- Processes can **directly** access (put / get) mem. location of a remote process
 - Remote process does not have to call a matching function

Parallel Architectures


- Distributed Memory
 - Each processing element (P) has its separate main memory block (M)
 - Data exchange is achieved through message passing over the network
 - Message passing could be either explicit (MPI) or implicit (BSP, PGAS)
 - Programs typically implemented as a set of OS entities with own (virtual) address spaces – *processes*
 - No shared variables
 - Usually no data races – but sometimes the opposite: missing data

Writing Programs for Distributed Memory Machines

- How do we write programs for such an architecture?
- **Classical approach: MPI (Message Passing Interface)** → HPC lecture
 - Defines an interface for sending messages between processors
 - Point-to-Point communication, Collective communication
- **More abstract approach: BSP (Bulk-Synchronous Parallel) programs**
 - Defines a generic structure of a parallel program (so called “supersteps”)
 - Program resp. algorithm has to be adapted to the BSP programming model
 - Provides a cost model to reason about performance
 - In practice: Often implemented on top of MPI

Literature


- Original paper on BSP: Valiant, L.G.. *A Bridging Model for Parallel Computation*. Communications of the ACM 33, 8 (August 1990), 103-111
– <http://doi.acm.org/10.1145/79173.79181>
- Parallel Algorithms Lecture (Rob Bisseling, Utrecht University)
– <http://www.staff.science.uu.nl/~bisse101/Education/PA/pa.html>
- Rob H. Bisseling: *Parallel Scientific Computation – A Structured Approach using BSP and MPI*, 2004, Oxford University Press
- Buurlage, J. W., Bannink, T., Bisseling, R. H.. *Bulk: A Modern C++ Interface for Bulk-Synchronous Parallel Programs*. In EuroPar 2018 (pp. 519-532). Springer
– https://doi.org/10.1007/978-3-319-96983-1_37




BSP Computer

BSP Model Components


BSP Computer (Distributed Memory Computer)



Programming Model (Algorithmic Framework)



Cost Model



BSP Computer

- BSP computer is a distributed memory computer
 - Consists of processors with own memory
 - Local memory accesses are fast, remote memory accesses slower
- Communication network treated as black box
 - Algorithm designer does not have to care about network details
- BSP algorithms are portable
 - Run on many different parallel computers (even on shared memory computers)

