



Concepts and Models of Parallel and Data-centric Programming

BSP II

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline


- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. **Bulk-Synchronous Parallelism**
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. Motivation
 - b. BSP Computer
 - c. BSP Programming Model
 - d. BSP Cost Model
 - e. Bulk Library

BSP Programming Model

BSP Model Components


BSP Computer

(Distributed Memory Computer)




Programming Model

(Algorithmic Framework)

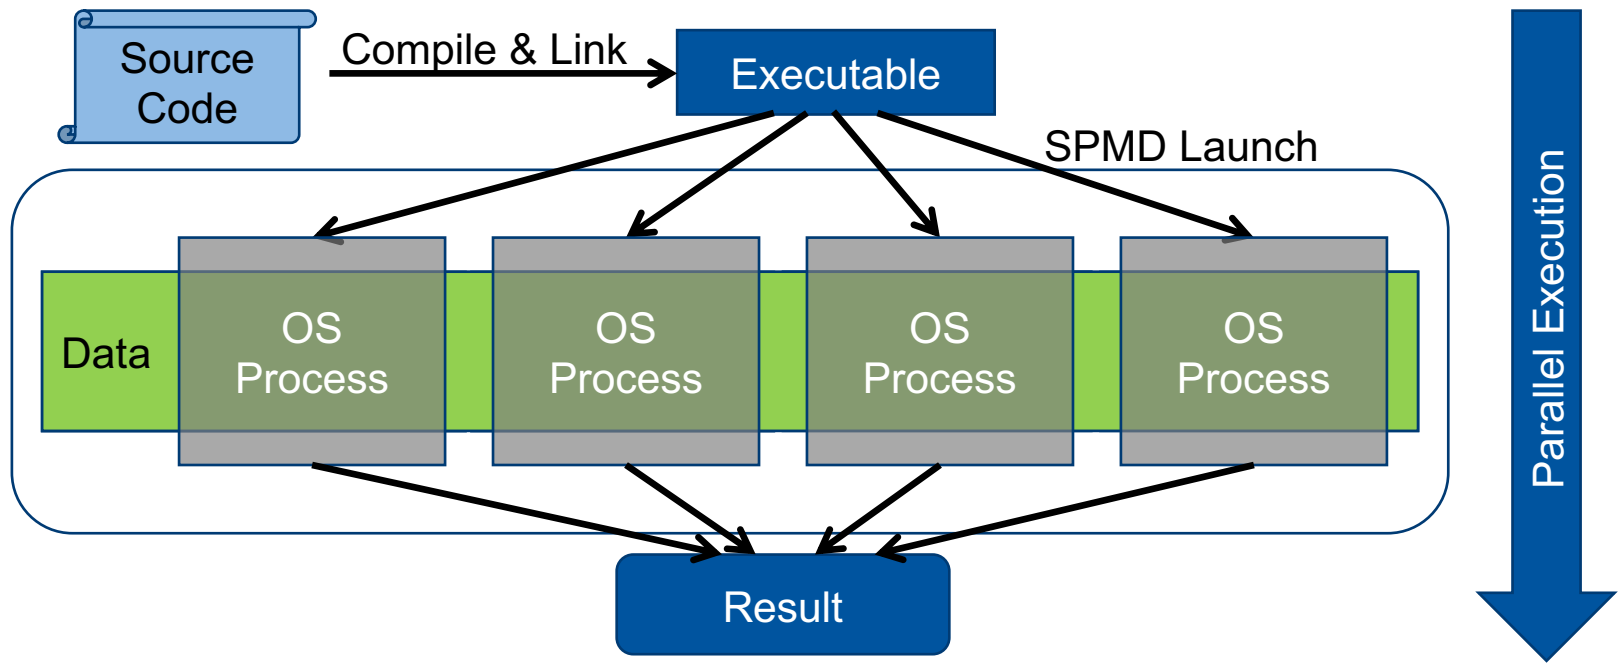


Cost Model




SPMD Model

- BSP programming model follows SPMD Program Lifecycle
 - Multiple processes run the same program with different data (Single Program Multiple Data)
- Program uses processor identity to differentiate processors (MPI: rank)



Programming Model: Supersteps

- BSP algorithm: Series of supersteps
- Superstep: Computation and communication
 - *Computation phase*: Perform local calculations with available data (e.g., FP operations)
 - *Communication phase*: Transfer data (e.g., results) between the different processors
- End of each superstep: Barrier
 - Each processor has to wait until all other processors have reached the barrier.
→ Bulk synchronization
 - Ensures that communication between processors has finished



Simple BSP Example: Parallel Inner Product Computation

- **Given:** Two vectors $x = (x_0, \dots, x_{n-1})^T$ and $y = (y_0, \dots, y_{n-1})^T$
- The *inner product* of x and y is defined as

$$\alpha = x^T y = \sum_{i=0}^{n-1} x_i y_i .$$

- **Goal:** Parallelize problem on a BSP computer with p processes
 - Result should be available at *all* processors in the end
- First step: Data distribution
- Second step: BSP algorithm design (supersteps)


First Step: Data Distribution

- **Given:** Vector $x = (x_0, \dots, x_{n-1})^T$, distribute to $p > 0$ processors
- **Distribution function:** $\text{distr}: \{0, \dots, n-1\} \rightarrow \{0, \dots, p-1\}$
- **Block distribution:** $i \mapsto \left\lfloor \frac{i}{b} \right\rfloor$, for $0 \leq i < n$ (x_i distributed to process $P(\left\lfloor \frac{i}{b} \right\rfloor)$)
where $b := \left\lceil \frac{n}{p} \right\rceil$ is the block size

Processors: $p = 4$

Vector length: $n = 16$

Block size: $b = 4$



- **Cyclic distribution:** $i \mapsto i \bmod p$, for $0 \leq i < n$

Processors: $p = 4$

Vector length: $n = 16$



Second Step: BSP Algorithm Design (1)


- Inner product: $\alpha = x^T y = \sum_{i=0}^{n-1} x_i y_i$
- Obvious: x_i and y_i should be on the same processor for each $0 \leq i < n$
 - Compute $x_i \cdot y_i$ locally
- Distribution should spread vector components **evenly** on the processors
 - Both block and cyclic distribution fulfill this requirement
- Choose cyclic distribution for x_i and y_i here

Second Step: BSP Algorithm Design (2)

Idea

- Superstep 1
 - Computation: Perform multiplications of locally available x_i and y_i , do a sum reduction on the products locally to get an intermediate result
 - Communication: Broadcast result to other processors
 - End of superstep: Barrier
- Superstep 2
 - Computation: Each processor performs a sum reduction on the received results
 - Communication: None, problem is solved. Result available on **all** processors.
 - End of superstep: Barrier (but not needed here)

Example for $n = 11$ and $p = 4$ (cyclic distribution)



BSP Algorithm: Parallel Inner Product (Cyclic Distribution)

Input: x, y : vector of length n



$distr(x) = distr(y) = d$ with $d(i) = i \bmod p$ for $0 \leq i < n$ (cyclic)

Output: $\alpha = x^T y$

every processor has its own rank
by this rank, each processor compute local product
these are distributed by processor

Algorithm for processor $s \in \{0, \dots, p - 1\}$:

$\alpha_s := 0;$	}	Comp. Phase	}	Superstep 0
for ($i := s; i < n; i += p$) do				
$\alpha_s := \alpha_s + x_i y_i;$ // compute local product				
for ($t := 0; t < p; t++$) do	}	Comm. Phase	}	
put α_s in $P(t);$ // broadcast to all processors t				
barrier();				
$\alpha := 0;$	}	Comp. Phase	}	Superstep 1 (omitted Comm.)
for ($t := 0; t < p; t++$) do				
$\alpha := \alpha + \alpha_t;$ // sum up local and received α values				