



Concepts and Models of Parallel and Data-centric Programming

BSP V (Bulk: Distributed Variables & Coarrays)

Lecture, Summer 2021

Simon Schwitanski

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. **Bulk-Synchronous Parallelism**
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. Motivation
 - b. BSP Computer
 - c. BSP Programming Model
 - d. BSP Cost Model
 - e. **Bulk Library**
 - a. Introduction
 - b. Data Distribution
 - c. **Distributed Variables**
 - d. **Coarrays**
 - e. **Further Features**

Bulk Library: Distributed Variables

Distributed Variables

- *Distributed variables* enable communication between processors
 - var object captures a distributed variable
- Has to be created in the same superstep by each processor
- Each processor has a *local value* of that variable and can access the concrete values on remote processors as so called *remote values*
- Accessing a remote value of a distributed variable `x`
 - `bulk::future<Type> y = x(t).get()` (read value of `x` at processor `t`)
 - `x(t) = value` (write value to `x` at processor `t`)

Question: Why is the `get()` operation encapsulated in a future?


- Because the communication (reading / writing a value) is only guaranteed to be completed after a `world.sync()` operation!

Distributed Variables – Example

```
1  env.spawn(env.available_processors(), [](auto& world) {
2      auto x = bulk::var<int>(world);
3      auto t = world.next_rank();
4      x(t) = 2 * world.rank();
5      world.sync();
6      // x now contains 2 * world.prev_rank() at the current rank
7      bulk::future<int> b = x(t).get(); // get value of next rank
8      world.sync();
9      // b.value() now contains two times the local rank ID
10 });
```

- Note: All communication via distributed variables is nonblocking!
 - `x(t).get();` and `x(t) = value` do not wait until a value has been read / or written

Example for 4 processes:



Translating BSP Algorithm to Bulk – Distributed Variables (1)

```
1  env.spawn(env.available_processors(),
2          [])(auto& world) {
3      auto s = world.rank();
4      auto p = world.active_processors();
5
6      // computation of local inner product
7      // assume: x and y cyclic partitioned arrays
8      double local_product = 0.0;
9      for (int i = 0; i < local_size; i++)
10         local_product += x.local(i) * y.local(i);
11
12     // create distributed variable for each rank
13     std::vector<bulk::var<double>> remote_products;
14     for (int t = 0; t < p; t++)
15         remote_products.
16             push_back(bulk::var<double>(world));
17
18     // put value to other processes
19     for (int t = 0; t < p; t++)
20         remote_products(t)[s] = local_product;
21
22     world.sync();
23
```

```
// compute local product
 $\alpha_s := 0;$ 
for (i := s; i < n; i += p) do
     $\alpha_s := \alpha_s + x_i y_i;$ 
```

```
// broadcast to all processors t
for (t := 0; t < p; t++) do
    put  $\alpha_s$  in P(t);
```

```
barrier();
```

```
// sum up received  $\alpha$  values
 $\alpha := 0;$ 
for (t := 0; t < p; t++) do
     $\alpha := \alpha + \alpha_t;$ 
```

local_size can be determined using the local_count() member function of the partitioning (not shown here)

Translating BSP Algorithm to Bulk – Distributed Variables (2)

```
24  // computation of global inner product
25  double inner_product = 0;
26  for (int t = 0; t < p; t++)
27      inner_product += remote_products[t];
28
29  // inner product available at each processor
30  // end of second superstep
31  });
```

```
// compute local product
 $\alpha_s := 0;$ 
for (i := s; i < n; i += p) do
     $\alpha_s := \alpha_s + x_i y_i;$ 

// broadcast to all processors t
for (t := 0; t < p; t++) do
    put  $\alpha_s$  in P(t);

barrier();

// sum up received  $\alpha$  values
 $\alpha := 0;$ 
for (t := 0; t < p; t++) do
     $\alpha := \alpha + \alpha_t;$ 
```

Bulk Library: Coarrays

Coarrays (1)

- **Problem:** Implementing a “broadcast” like shown before with distributed variables is ugly
- What we want to achieve: Put our local value to a defined place on the other processors
- Idea: Allocate an array of size p on each processor
 - Each processor s puts its local value at the s -th place in the array of all other processors


Coarrays (2)

- Approach: Coarray
 - Model distributed data as 2D array
 - First dimension: Processor
 - Second dimension: Chosen processor's 1D array (according to first dimension)
- Examples:
 - `coarray(t)[41] = value`
 - Put value to the 42-th element in the 1D array of processor t
 - `coarray(t)[{start,end}] = {value1, value2, ...}`
 - Put value1, value2, ... to the elements start, start + 1, ..., end - 1 at processor t (modifying *slices* of arrays)

Coarrays (3)

```
1  env.spawn(env.available_processors(), [](auto& world) {
2      auto xs = bulk::coarray<int>(world, 4); // coarray of local size 4
3
4      // each processor sets its first local element to 1
5      xs[0] = 1;
6      auto t = world.next_rank();
7      // set second element of processor t's array to rank ID
8      xs(t)[1] = world.rank();
9      // set third and fourth element of processor t's array to 8 and 9
10     xs(t)[{2,4}] = {8, 9};
11 });
```

Example for 4 processes:



Translating BSP Algorithm to Bulk – Coarrays (1)

```
1 env.spawn(env.available_processors(),
2           [](auto& world) {
3     auto s = world.rank();
4     auto p = world.active_processors();
5
6     // computation of local inner product
7     // assume: x and y cyclic partitioned arrays
8     double local_product = 0.0;
9     for (int i = 0; i < local_size; i++)
10        local_product += x.local(i) * y.local(i)
11
12    // broadcast to all processors t
13    auto remote_products =
14        bulk::coarray<double>(world, p);
15    for (int t = 0; t < p; t++)
16        remote_products(t)[s] = local_product;
17
18    world.sync(); // end of first superstep
19
```

```
// compute local product
 $\alpha_s := 0;$ 
for (i := s; i < n; i += p) do
     $\alpha_s := \alpha_s + x_i y_i;$ 
```

```
// broadcast to all processors t
for (t := 0; t < p; t++) do
    put  $\alpha_s$  in P(t);
```

```
barrier();
```

```
// sum up received  $\alpha$  values
 $\alpha := 0;$ 
for (t := 0; t < p; t++) do
     $\alpha := \alpha + \alpha_t;$ 
```

Translating BSP Algorithm to Bulk – Coarrays (2)

```
20 // computation of global inner product
21 double inner_product = 0;
22 for (int t = 0; t < p; t++)
23     inner_product += remote_products[t];
24
25 // inner product available at each processor
26 // end of second superstep
27 });
```

```
// compute local product
 $\alpha_s := 0;$ 
for (i := s; i < n; i += p) do
     $\alpha_s := \alpha_s + x_i y_i;$ 

// broadcast to all processors t
for (t := 0; t < p; t++) do
    put  $\alpha_s$  in P(t);

barrier();

// sum up received  $\alpha$  values
 $\alpha := 0;$ 
for (t := 0; t < p; t++) do
     $\alpha := \alpha + \alpha_t;$ 
```

Data Exchange in Bulk: Summary

Object	Image	Description	Code
var	local	set	<code>x = 5</code>
		use	<code>auto y = x + 3</code>
	remote	put	<code>x(t) = 5</code>
		get	<code>auto y = x(t).get()</code>
coarray	local	set	<code>xs[idx] = 5</code>
		use	<code>auto y = xs[idx] + 3</code>
	remote	put	<code>xs(t)[idx] = 5</code>
		get	<code>auto y = xs(t)[idx].get()</code>
		put slice*	<code>xs(t)[{start, end}] = {values...}</code>
		get slice*	<code>auto ys = xs(t)[{start, end}].get()</code>


Source: Buurlage, J. W., Bannink, T., Bisseling, R. H.. *Bulk: A Modern C++ Interface for Bulk-Synchronous Parallel Programs*.

*: Slices are represented as `std::vector` containers


Bulk Library: Further Features

Further Features (1)

- Common communication patterns also available, e.g.:
 - `broadcast()`



- `gather_all()` (what we did manually in the inner product algorithm)



Further Features (2)

- Mailbox communication: queue object enables message passing
 - Similar to MPI message passing
- Special-purpose objects for vectors, matrices, ...
- Pre-implemented algorithms (inner product, sorting, LU decomposition)
- Further reading:
 - <https://jwbuurlage.github.io/Bulk/>
 - Buurlage, J. W., Bannink, T., Bisseling, R. H.. *Bulk: A Modern C++ Interface for Bulk-Synchronous Parallel Programs*. In EuroPar 2018 (pp. 519-532). Springer
 - https://doi.org/10.1007/978-3-319-96983-1_37

What you have learnt

- Bulk: C++ library implementing the BSP model
 - Supporting different backends: MPI, C++ threads, coprocessors
- Data distribution objects
- Communication mechanisms
 - Distributed variables
 - Coarrays
- Implementation of inner product computation in Bulk
- Other supported communication patterns