



Concepts and Models of Parallel and Data-centric Programming

Distributed Shared Memory

Lecture, Summer 2020

Dr. Christian Terboven <terboven@itc.rwth-aachen.de>

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory**
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. MapReduce
 - 10. Apache Spark
- a. PGAS Foundations
 - b. DASH Overview
 - c. Distributed Data Structured
 - d. DASH Algorithms
 - e. Tasking

Distributed Data Structures

Distributed Data Structures

- DASH offers distributed data structures
 - Support for flexible data distribution schemes
 - Example: `dash::Array<T>`

```
dash::Array<int> arr(100);
```

```
if( dash::myid()==0 ) {  
    for( auto i=0; i<arr.size(); i++ )  
        arr[i]=i;  
}
```

```
arr.barrier();  
if(dash::myid()==1 ) {  
    for( auto el: arr )  
        cout<<(int)el<<" ";  
    cout<<endl;  
}
```

DASH global array of 100 integers,
distributed over all units,
default distribution is BLOCKED

Unit 0 writes to the array using the
global index `i`. Operator `[]` is
overloaded for the `dash::Array`.

Unit 1 executes a range-
based for loop over the
DASH array

```
$ mpirun -n 4 ./array
```

```
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19  
20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36  
37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53  
54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70  
71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87  
88 89 90 91 92 93 94 95 96 97 98 99
```

Accessing Local Data

- Access to the local portion of the data is exposed through a local-view proxy object (.local)

```
dash::Array<int> arr(100);  
  
for( auto i=0; i<arr.lsize(); i++ )  
    arr.local[i]=dash::myid();  
  
arr.barrier();  
if(dash::myid()==dash::size()-1 ) {  
    for( auto el: arr )  
        cout<<(int)el<<" ";  
    cout<<endl;  
}
```

.lsize() is short hand for **.local.size()** and returns the number of local elements

.local is a *proxy object* that represents the part of the data that is local to a unit.

```
$ mpirun -n 4 ./array
```

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
2 2 2 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3  
3 3 3 3
```

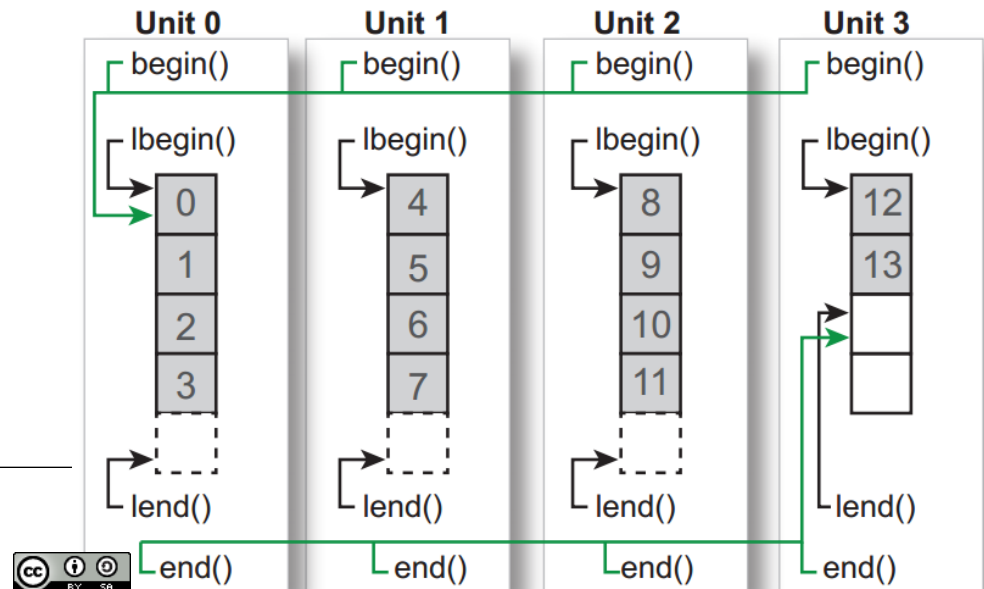
Global-view and Local-view in DASH

- DASH supports both *global-view* and *local-view* semantics

	Global-view	Local-view	LV shorthand
range begin	<code>arr.begin()</code>	<code>arr.local.begin()</code>	<code>arr.lbegin()</code>
range end	<code>arr.end()</code>	<code>arr.local.end()</code>	<code>arr.lend()</code>
# elements	<code>arr.size()</code>	<code>arr.local.size()</code>	<code>arr.lsize()</code>
element access	<code>arr[glob_idx]</code>	<code>arr.local[loc_idx]</code>	

- Example

- `dash::Array` with 14 elements, distributed over 4 units
- default distribution: BLOCKED
- $\text{Blocksize} = \text{ceil}(14/4)=4$



Efficient Local Access (1)

- Several options for access to local data

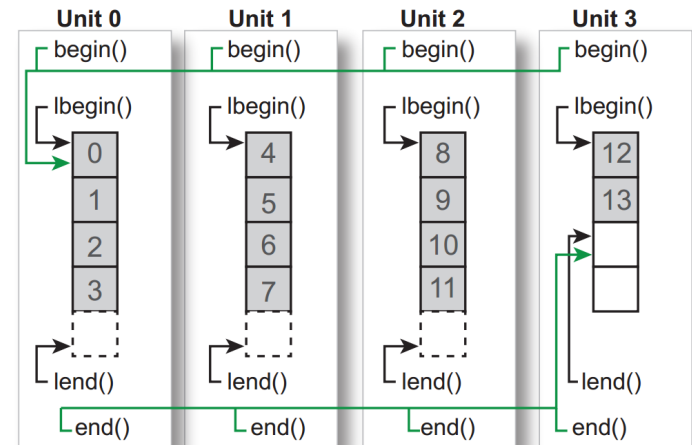
```
dash::Array<int> arr(1000);

// get raw pointer to local mem.
int *p1 = arr.local.begin();
int *p2 = arr.lbegin(); //p1==p2

// access via local index
arr.local[22] = 33;

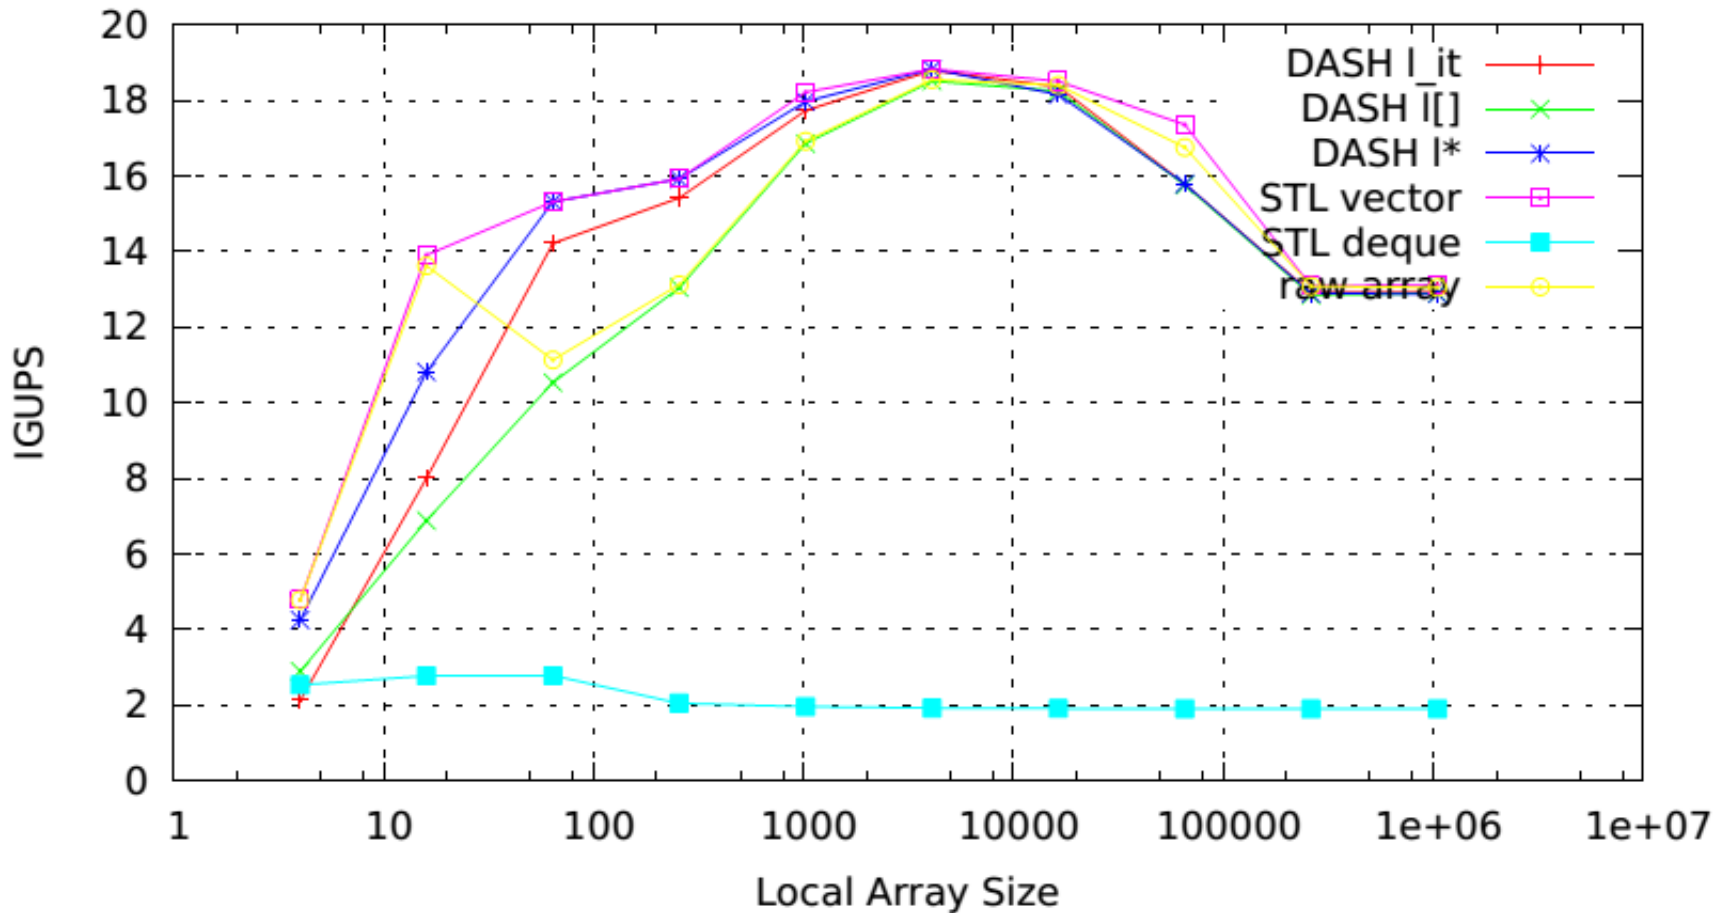
// range-based for loop
for( auto el : arr.local )
    cout<<el<<" ";

// access using local iterators
for( auto it=arr.lbegin(); it!=arr.lend(); ++it ) {
    (*it) = foo(...);
}
```



Efficient Local Access (2)

- IGUPs Benchmark: independent parallel updates



Using STL Algorithms

- STL algorithms can be used with DASH containers
 - Both on the local view and the global view

```
#include <libdash.h>
```

```
int main(int argc, char* argv[])  
{
```

```
    dash::init(&argc, &argv);
```

```
    dash::Array<int> a(1000);
```

```
    if( dash::myid()==0 ) {  
        // global iterators and std. algorithms  
        std::sort(a.begin(), a.end());  
    }
```

```
    // local access using local iterators  
    std::fill(a.lbegin(), a.lend(), 23+dash::myid());
```

```
    dash::finalize();
```

```
}
```

Collective constructor,
all units involved

STL algorithms work
with DASH global
iterators

STL algorithms work with
DASH local iterators


Data Distribution Patterns

- The data distribution pattern is configurable

```
dash::Array<int> arr1(20); // default: BLOCKED  
  
dash::Array<int> arr2(20, dash::BLOCKED)  
dash::Array<int> arr3(20, dash::CYCLIC)  
dash::Array<int> arr4(20, dash::BLOCKCYCLIC(3))
```


- Assume 4 units

arr1, arr2




BLOCKED

arr3



CYCLIC

arr4



BLOCKCYCLIC(3)

Accessing Local Data – Cyclic Distribution

- The previous example with a cyclic distribution:

```
// this is the only changed line
dash::Array<int> arr(100, dash::CYCLIC);

for( auto i=0; i<arr.lsize(); i++ )
    arr.local[i]=dash::myid();

arr.barrier();
if(dash::myid()==dash::size()-1 ) {
    for( auto el: arr )
        cout<<(int)el<<" ";
    cout<<endl;
}
```

```
$ mpirun -n 4 ./array
```

```
0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3 0
1 2 3 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3 0 1
2 3 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2
3 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3
0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3
```

DASH Distributed Data Structures Overview

Container	Description	Data distribution
Array <T>	1D Array	static, configurable
NArray <T, N>	N-dim. Array	static, configurable
Shared <T>	Shared scalar	fixed (at 0)
Directory ^(*) <T>	Variable-size, locally indexed array	manual


(*) Under development

Multidimensional Data Distribution (1)


- dash::Pattern<N> specifies N-dim data distribution
 - Blocked, cyclic, and block-cyclic in multiple dimensions

Pattern<2>(20, 15)

Extent in first and second dimension



(BLOCKED,
NONE)



(NONE,
BLOCKCYCLIC(2))



(BLOCKED,
BLOCKCYCLIC(3))


■ Unit 0
■ Unit 1
■ Unit 2
■ Unit 3

Distribution in first and second dimension

Multidimensional Data Distribution (2)


- Tiled data distribution and tile-shifted distribution

TilePattern<2>(20, 15)



(TILE(2), TILE(5))

ShiftTilePattern<2>(32, 24)




(TILE(4), TILE(3))



Multidimensional Data Distribution (3)


- Row-major and column-major storage

Pattern<2, COL_MAJOR>(20, 15)



(BLOCKED, NONE)

TilePattern<2, COL_MAJOR>(20, 15)



(TILE(5), TILE(5))

Unit 0
Unit 1
Unit 2
Unit 3

The N-Dimensional Array

- `dash::NArray` (`dash::Matrix`) offers a distributed multidimensional array abstraction
 - Dimension is a template parameter
 - Element access using coordinates or linear index
 - Support for custom index types
 - Support for row-major and column-major storage

```
dash::NArray<int, 2> mat(40, 30); // 1200 elements


int a = mat(i,j);    // Fortran style access
int b = mat[i][j];   // chained subscripts

auto loc = mat.local;

int c = mat.local[i][j];
int d = *(mat.local.begin()); // local iterator
```


DASH NArray Global View and Local View

- Local view works similar to 1D array



Global View

(0,0) 0 a	(0,1) 1 b	(0,2) 2 c	(0,3) 3 d
(1,0) 4 e	(1,1) 5 f	(1,2) 6 g	(1,3) 7 h
(2,0) 8 i	(2,1) 9 j	(2,2) 10 k	(2,3) 11 l
(3,0) 12 m	(3,1) 13 n	(3,2) 14 o	(3,3) 15 p
(4,0) 16 q	(4,1) 17 r	(4,2) 18 s	(4,3) 19 t
(5,0) 20 u	(5,1) 21 v	(5,2) 22 w	(5,3) 23 x
(6,0) 24 y	(6,1) 25 z	(6,2) 26 A	(6,3) 27 B

```
dash::NArray<char, 2> mat(7, 4);  
cout << mat(2, 1) << endl; // prints 'j'  
  
if(dash::myid()==0) {  
    cout << mat.local(2, 1) << endl; // prints 'z'  
}
```

Local View (Unit 0)

(0,0) 0 a	(0,1) 1 b	(0,2) 2 c	(0,3) 3 d
(1,0) 4 m	(1,1) 5 n	(1,2) 6 o	(1,3) 7 p
(2,0) 8 y	(2,1) 9 z	(2,2) 10 A	(2,3) 11 B

Local View (Unit 1)

(0,0) 0 e	(0,1) 1 f	(0,2) 2 g	(0,3) 3 h
(1,0) 4 q	(1,1) 5 r	(1,2) 6 s	(1,3) 7 t

Local View (Unit 2)

(0,0) 0 i	(0,1) 1 j	(0,2) 2 k	(0,3) 3 l
(1,0) 4 u	(1,1) 5 v	(1,2) 6 w	(1,3) 7 x