



# Concepts and Models of Parallel and Data-centric Programming

Parallel Algorithms IV

Lecture, Summer 2020

Dr. Christian Terboven <[terboven@itc.rwth-aachen.de](mailto:terboven@itc.rwth-aachen.de)>

# Outline


---

- 0. Organization
- 1. Foundations
- 2. Shared Memory
- 3. GPU Programming
- 4. Bulk-Synchronous Parallelism
- 5. Message Passing
- 6. Distributed Shared Memory
- 7. Parallel Algorithms**
  - a. Berkeley DWARFS
  - b. Dense Linear Algebra
  - c. Sparse Linear Algebra
  - d. Monte Carlo Methods
  - e. Graph Traversal
- 8. Parallel I/O
- 9. MapReduce
- 10. Apache Spark

# Monte Carlo

---

- $\frac{\pi}{4}$  = ratio of
  - area of the circle
  - area of the square



- Monte Carlo method
  - Statistical simulation
  - The ratio of the number of points that fall inside **A** to the total number of points tried (all within **R**) is equal to the ratio of the two areas (or volume in 3d)
    - Randomly throw darts inside a 2x2 square area
    - Count darts that hit the radius 1 circle
    - $PI = 4 * \text{hits}/\text{total}$


- What is the intuitive approach to parallelize this problem?
- Nearly embarrassingly parallel solver
  - Parallel generation of random numbers
  - Parallel counting of hits
  - Linear collection of hit count at rank 0 (by design)
  - Final calculation and printout at rank 0
- Scaling expectation
  - Scales linear until sequential part starts to dominate (Amdahl's law)

- Perfect network, i.e., 0 latency and  $\infty$  bandwidth

- Eventually, the applications runs out of work

- Increasing the number of iterations pushes the scalability limit

## Scaling with Perfect Network (AMD Opteron CPU, 0 Latency, $\infty$ Bandwidth)



C. Engelmann. Scaling To A Million Cores And Beyond: A Basic Understanding Of The Challenges Ahead On The Road To Exascale.