



Concepts and Models of Parallel and Data-centric Programming

MapReduce – Hadoop Distributed File System

Lecture, Summer 2020

Simon Schwitanski
Dr. Christian Terboven

Outline

0. Organization
 1. Foundations
 2. Shared Memory
 3. GPU Programming
 4. Bulk-Synchronous Parallelism
 5. Message Passing
 6. Distributed Shared Memory
 7. Parallel Algorithms
 8. Parallel I/O
 - 9. MapReduce**
 10. Apache Spark
- a. MapReduce Programming Model
 - b. Parallelizing MapReduce
 - c. Hadoop Ecosystem
 - d. Hadoop Distributed File System**
 - e. Yet Another Resource Negotiator
 - f. Comparison to Other Approaches
 - g. MapReduce Design Patterns

Distributed File System

- Datasets to analyze much larger than storage capacity of single machine
- Idea: Partition data across several machines
- *Distributed File System*: Managing files across network of machines
- Hadoop Distributed File System (HDFS)
 - De-facto standard storage system in Apache Hadoop

HDFS Design Goals (1)

Large Datasets

- Applications running on HDFS have large datasets
- Typical file size: Gigabytes to terabytes

Streaming Data Access

- Write-once-read-many access pattern
- Dataset generated or copied once, different analyses on same dataset
- Time to read complete dataset more important than low latency access

Commodity Hardware

- No special purpose hardware (interconnect, CPU, ...) required
- “Cheap” components

HDFS Design Goals (2)

Data Locality

- “Moving computation is cheaper than moving data”
- Especially for huge data sets
- Perform computation (physically) close to data
- Minimizes network congestion, increases throughput

Fault Tolerance


- 100 – 10,000 machines, each storing part of the data
- Probability of failures high
- Quick detection and recovery of faults

Portability

- Support of heterogeneous hardware and software platforms

HDFS Blocks

- HDFS: Logical abstraction of underlying physical file system (ext4, ...)
- Each file split into equally-sized chunks called *blocks*
- Blocks treated as independent units in file system
- Block size larger (128 MB) than in physical file systems (1 – 8 KB)
 - File smaller than single block does not occupy full block




Benefits of Block Abstraction

- Files can be larger than storage capacity of single machine
 - Store blocks of file on different machines
- Simplified storage management (mostly equally-sized blocks)
- Fault tolerance easier to implement
 - Replicate blocks on different machines
 - Lost block (machine failure, corruption) recovered using replication
 - *Replication factor*: Number of replicated blocks (default: 3)

NameNode and DataNode

- Master / Worker architecture
- NameNode: (Single) Master
 - Manages file system namespace and metadata (e.g., timestamp, length)
 - Provides access to clients
- DataNode: Worker
 - Manages local storage attached to node (as HDFS blocks)
 - Periodically sends heartbeats to NameNode with list of stored blocks
- Client accesses data by communicating with NameNode and DataNodes



HDFS Architecture

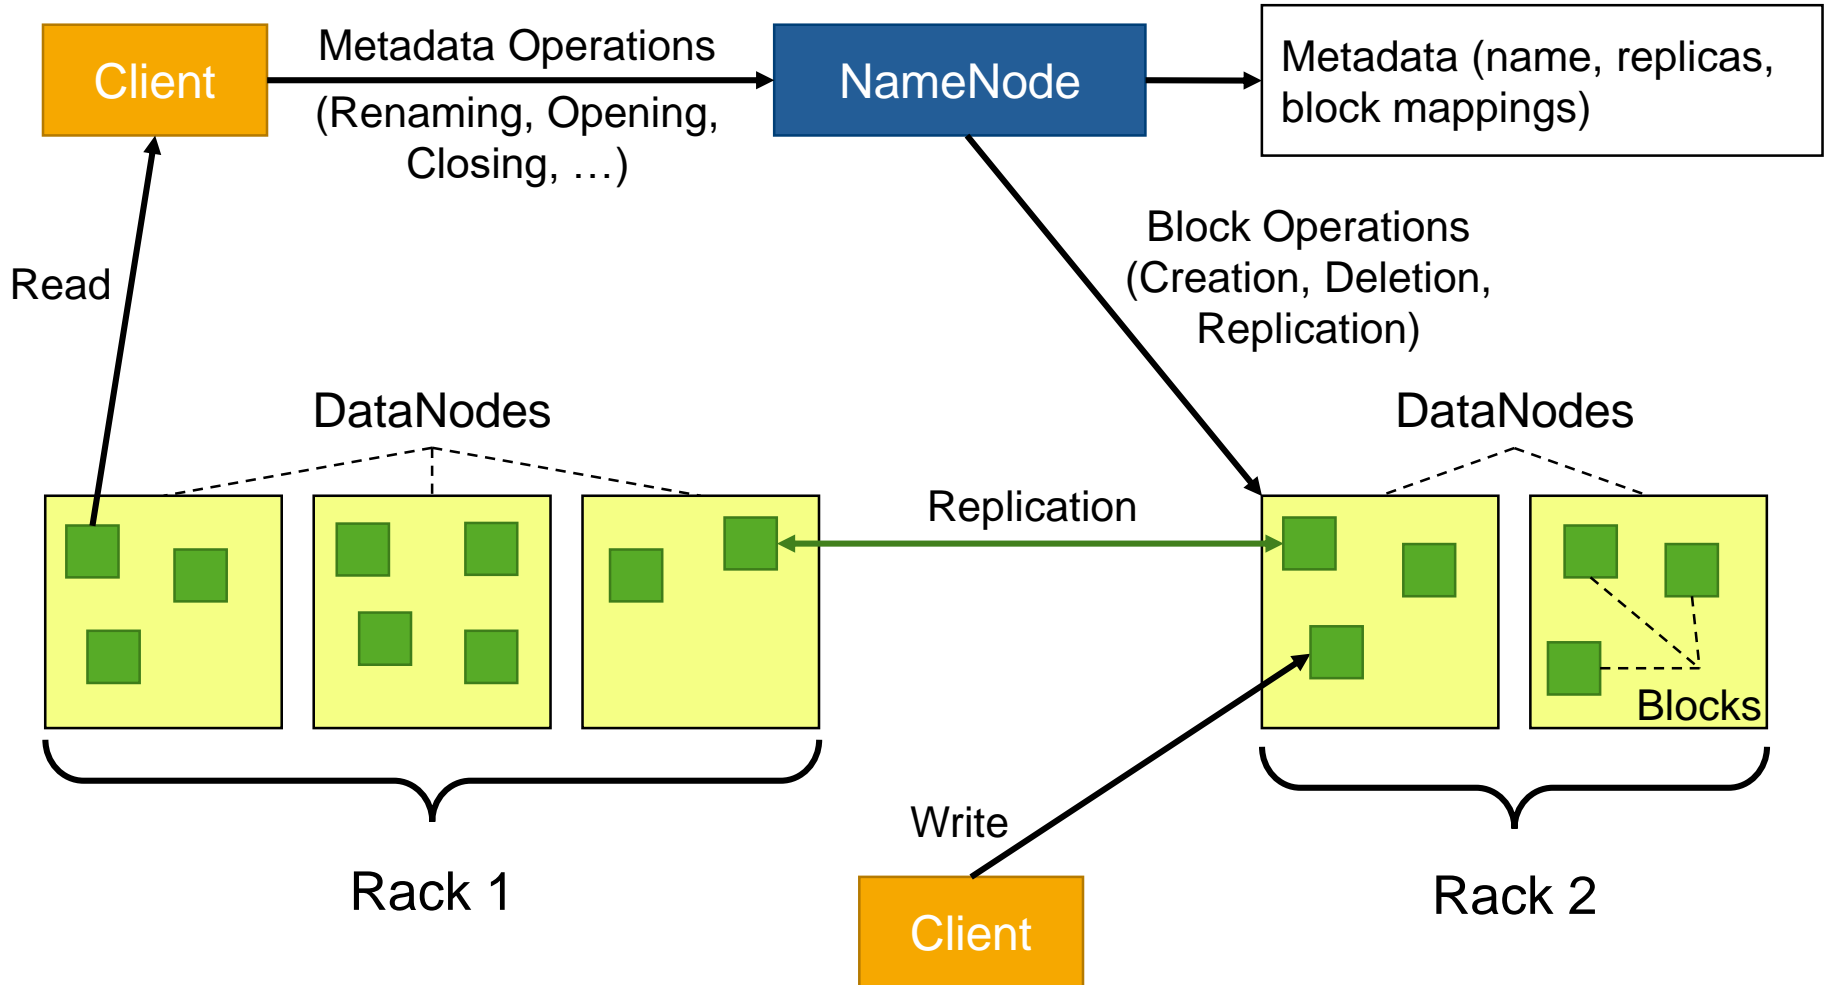
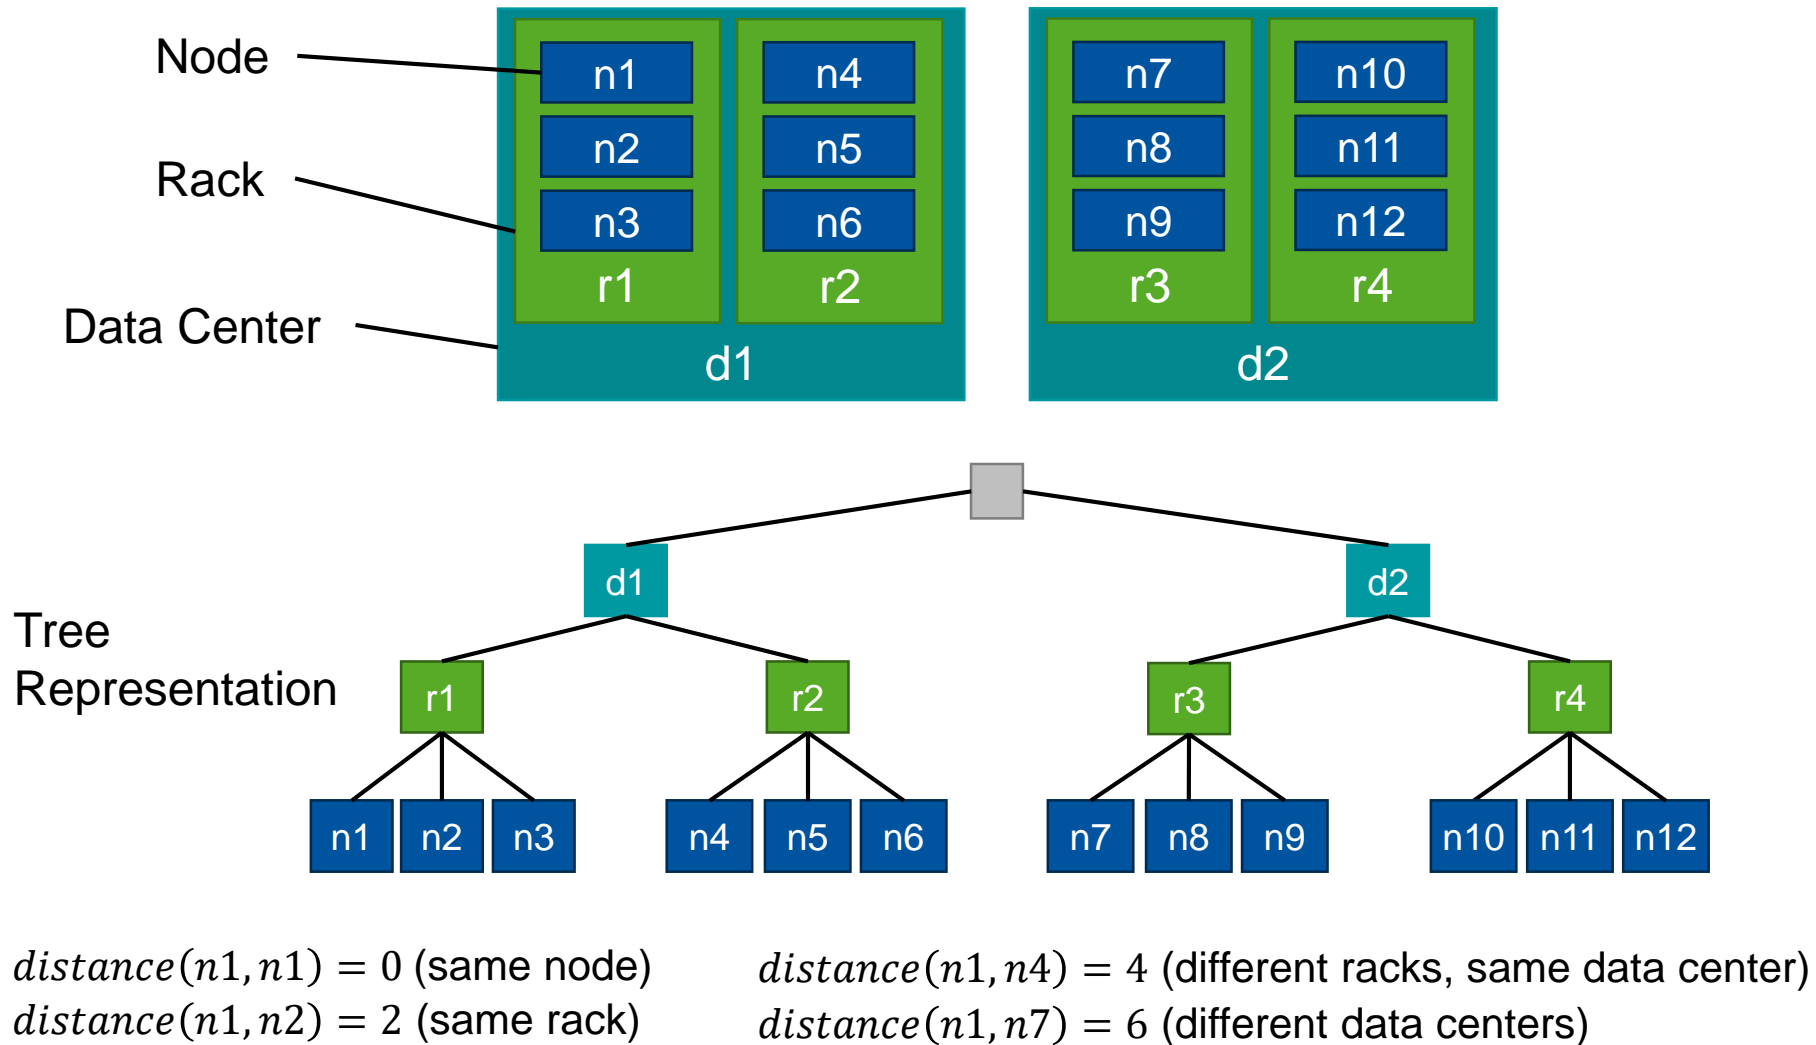


Illustration adapted from <https://hadoop.apache.org/docs/r3.1.0/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>

Network Topology in Hadoop


- DataNodes with requested blocks should be “close” to client node
 - Client node: Node that requests data blocks (for computation)
- Network bandwidth: Limiting factor when processing huge datasets
- Possible distance measure for two nodes: Bandwidth
 - Problem 1: Measuring bandwidth in a cluster is difficult
 - Problem 2: Number of pairs is in $O(n^2)$ for n nodes
- Distance measure in Hadoop
 - Network represented as tree
 - Distance between two nodes: Minimum distance between nodes in tree
 - DataNode with smallest distance to client node is chosen
 - Hadoop has to be made aware of network structure

Network Topology – Example



HDFS Architecture – Revisited

- NameNode is arbitrator for **all** metadata
 - Client asks NameNode before reading from or writing to any DataNode
 - User data written **directly** to DataNode (avoid overload of NameNode)
- Any problems with this approach?




NameNode Failure

- NameNode is ***Single Point of Failure (SPOF)***
- If machine running NameNode fails
 - All data lost, no way of reconstruction
- **Solution:** Backup all relevant data on external file systems
- On failure: Restore file system state from backup data
 - Read file system image backup
 - Apply (potentially buffered) metadata modifications
 - Wait for DataNodes reporting blocks (block mapping lost on NameNode failure)
- **However:** SPOF problem remains
 - Recovery can take about 30 minutes depending on data size

HDFS High Availability (Hadoop 2.0 onwards)

- Add a standby NameNode (active-standby configuration)
- On failure of primary node: Standby node takes over



DataNode Failure

- DataNode periodically sends heartbeats to NameNode
 - Contains report of stored blocks
- DataNode failure detection: Absence of heartbeat messages
- Default timeout: 10 minutes
- NameNode marks DataNode without heartbeats as dead
 - No further requests sent to this node
 - If replication factor below specified value: Re-replication on other DataNode
- Not covered in this lecture: Data integrity