



# Concepts and Models of Parallel and Data-centric Programming

MapReduce – Yet Another Resource Negotiator

Lecture, Summer 2020

Simon Schwitanski  
Dr. Christian Terboven

# Outline

---

0. Organization
  1. Foundations
  2. Shared Memory
  3. GPU Programming
  4. Bulk-Synchronous Parallelism
  5. Message Passing
  6. Distributed Shared Memory
  7. Parallel Algorithms
  8. Parallel I/O
  9. **MapReduce**
  10. Apache Spark
- a. MapReduce Programming Model
  - b. Parallelizing MapReduce
  - c. Hadoop Ecosystem
  - d. Hadoop Distributed File System
  - e. **Yet Another Resource Negotiator**
  - f. Comparison to Other Approaches
  - g. MapReduce Design Patterns

# Yet Another Resource Negotiator

---

- Cluster resource management system of Hadoop
- Introduced in Hadoop 2, encapsulates resource management
- General enough to support other computing paradigms
- YARN API: Request and run applications on cluster resources
  - Typically, user gets not directly in contact with YARN API
  - Higher-layer frameworks (e.g., MapReduce) used

**Application**

MapReduce

**Compute**


Yet Another Resource Negotiator (YARN)

**Storage**

Hadoop Distributed Filesystem (HDFS)


# Core Components (1)

- Master / Worker architecture
- Resource Manager (RM)
  - Managing resources across whole cluster
  - One per cluster
  - Schedules jobs to NodeManager nodes
- Node Manager (NM)
  - Runs on each worker node
  - Responsible for running and monitoring *containers*
- Container
  - Represents allocated set of resources (memory, CPU, ...) on a node
  - Application-specific task runs in container



## Core Components (2)

- Application master
  - One instance per application
  - Started by RM, runs in a container on arbitrary NM
  - Negotiates further resources (i.e., further containers) with RM
  - Executes and monitors map and reduce tasks in containers on different NMs
  - Job coordination and tracking
- Client
  - Submits job to RM
  - RM creates application master on arbitrary NM, starts computation
  - Polls application master for job progress



# YARN Architecture

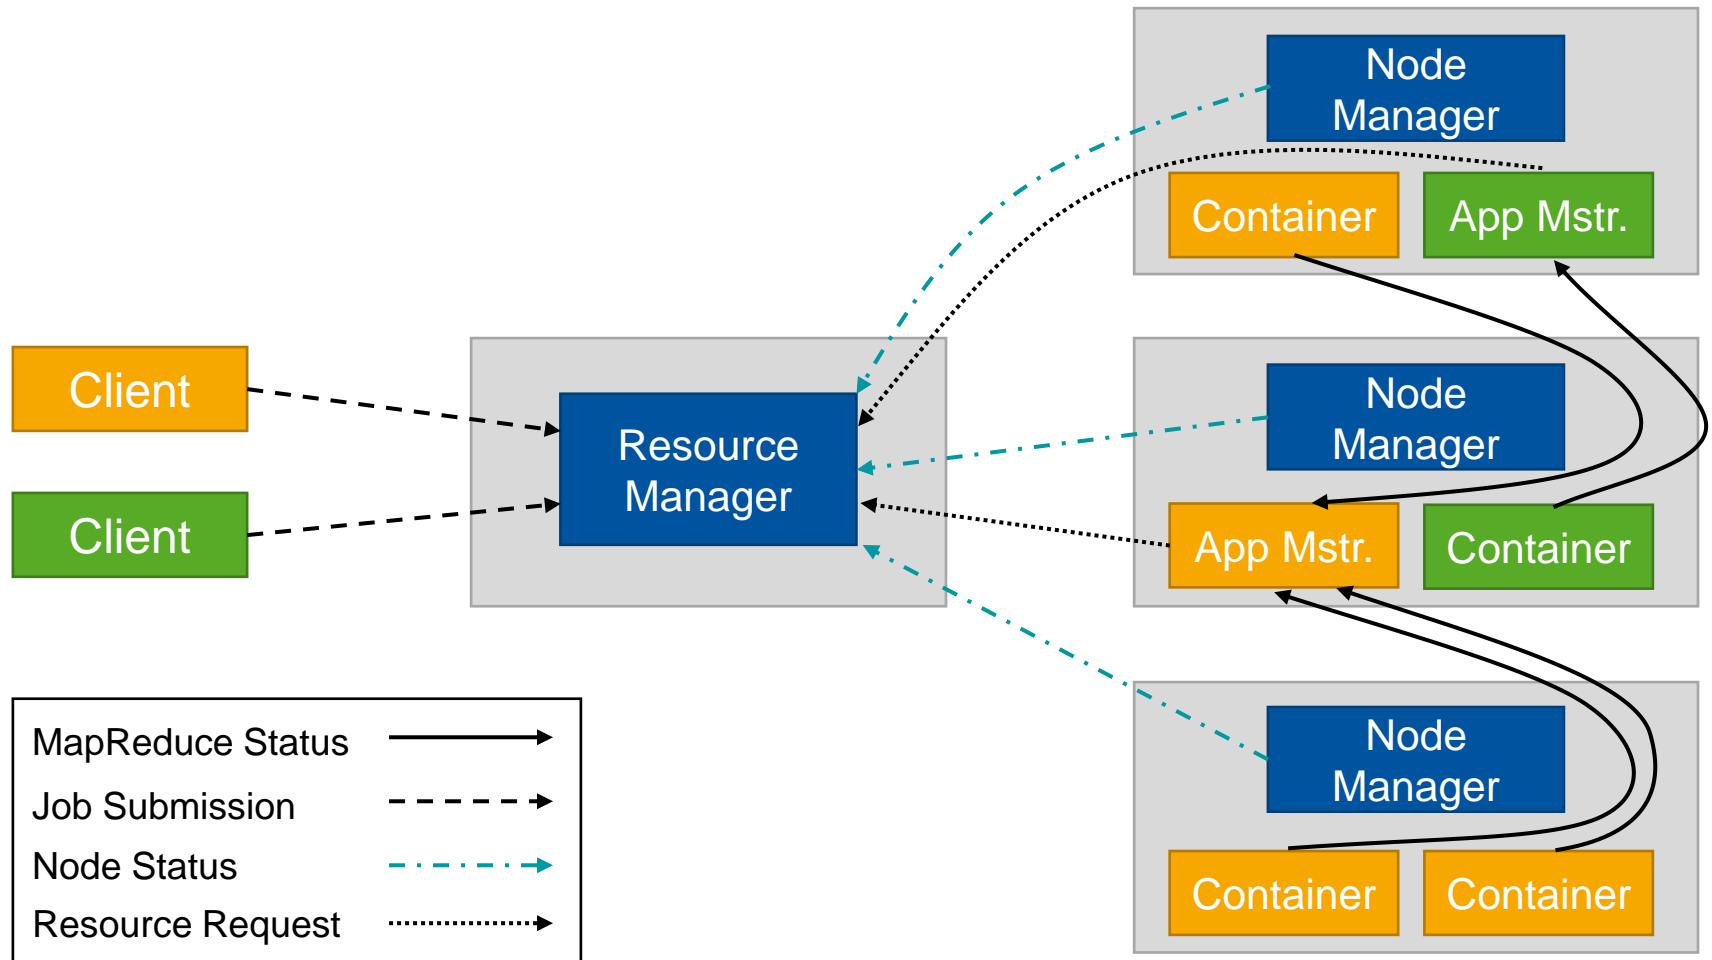


Illustration adapted from <https://hadoop.apache.org/docs/r3.1.0/hadoop-yarn/hadoop-yarn-site/YARN.html>

# YARN and HDFS

---

- Both YARN and HDFS: Master / Worker architecture
- Typical cluster setup: Two machines as masters and rest workers
  - HDFS NameNode on one machine (master)
  - ResourceManager on another machine (master)
  - Other machines: HDFS DataNode and NodeManager simultaneously (workers)
- Storage handling done by HDFS
- Computation handling done by YARN

# Task Assignments – Data Locality (1)

---

- Performance bottleneck for data-intensive task: Network bandwidth
- Data management concept of HDFS: **Data Locality**
- Container request in YARN can have locality constraints
  - Request to run task on a certain node
- Application master knows from NameNode which data lies on which node
  - Include locality constraint in container request
- Map task: Processing HDFS blocks locally
  - Ideally: Container runs on a node storing a replica of accessed HDFS blocks



# Task Assignments – Data Locality (2)

---

- Locality levels
  - Data-local (optimal, task runs on same node which stores needed data)
  - Rack-local (same rack, but not same node)
- Data locality only possible for map tasks
  - Reduce tasks collect data from different machines, no data locality
- **Goal:** Process most input data (of map tasks) locally → Less network bandwidth needed

# Task Granularity

---

- Map phase:  $M$  map tasks
  - Reduce phase:  $R$  reduce tasks
  - Ideally:  $M$  and  $R$  much larger than the number of worker nodes
    - Better load balancing
    - Faster recovery in case of a failure
  - Too large values for  $M$  and  $R$  can lead to significant overhead at ResourceManager
- 
- In practice:  $M = 200,000$  and  $R = 5,000$  with 2,000 worker machines, each map task processes about 16 MB to 64 MB (Google, 2008)


# Fault Tolerance

---

- MapReduce programs run on large number of machines
- Implementation must tolerate and handle process crashes, machine failures, ...
- Reasons for failures
  - Bad user codes, hardware faults, network connection issues, OS crash, ...
- Four failure types
  - Task failure: Reschedule task or abort job execution
  - Application master failure: Start new application master on other NM, job recovery required
  - **NodeManager failure**
  - **ResourceManager failure**


# Fault Tolerance – Node Manager Failure

- Periodic heartbeats sent from NM to RM
- On NM failure: Absence of heartbeats
- RM detects problem, potentially recovers corresponding tasks and application master
- Additionally: Completed map tasks have to be rerun for incomplete jobs
  - Intermediate output on the failed NM's local file system might be inaccessible




# Fault Tolerance – Resource Manager Failure

- Resource manager is single point of failure (SPOF)
  - No jobs or task containers can be started without it
- Similar to HDFS NameNode: Run pair of resource managers to resolve SPOF
  - Active-standby configuration
  - State is stored in shared storage
- On failure of active resource manager
  - Standby resource manager can recover state from storage and take over




# Speculative Execution (1)

- Critical for total execution time: Straggler
  - Machine taking unusually long time to complete one of the last map or reduce tasks
- Reasons for stragglers
  - Bad disk with slow read performance
  - Other tasks on the same machine
  - Bugs in machine configuration
- Probability of a straggler high for thousands of machines

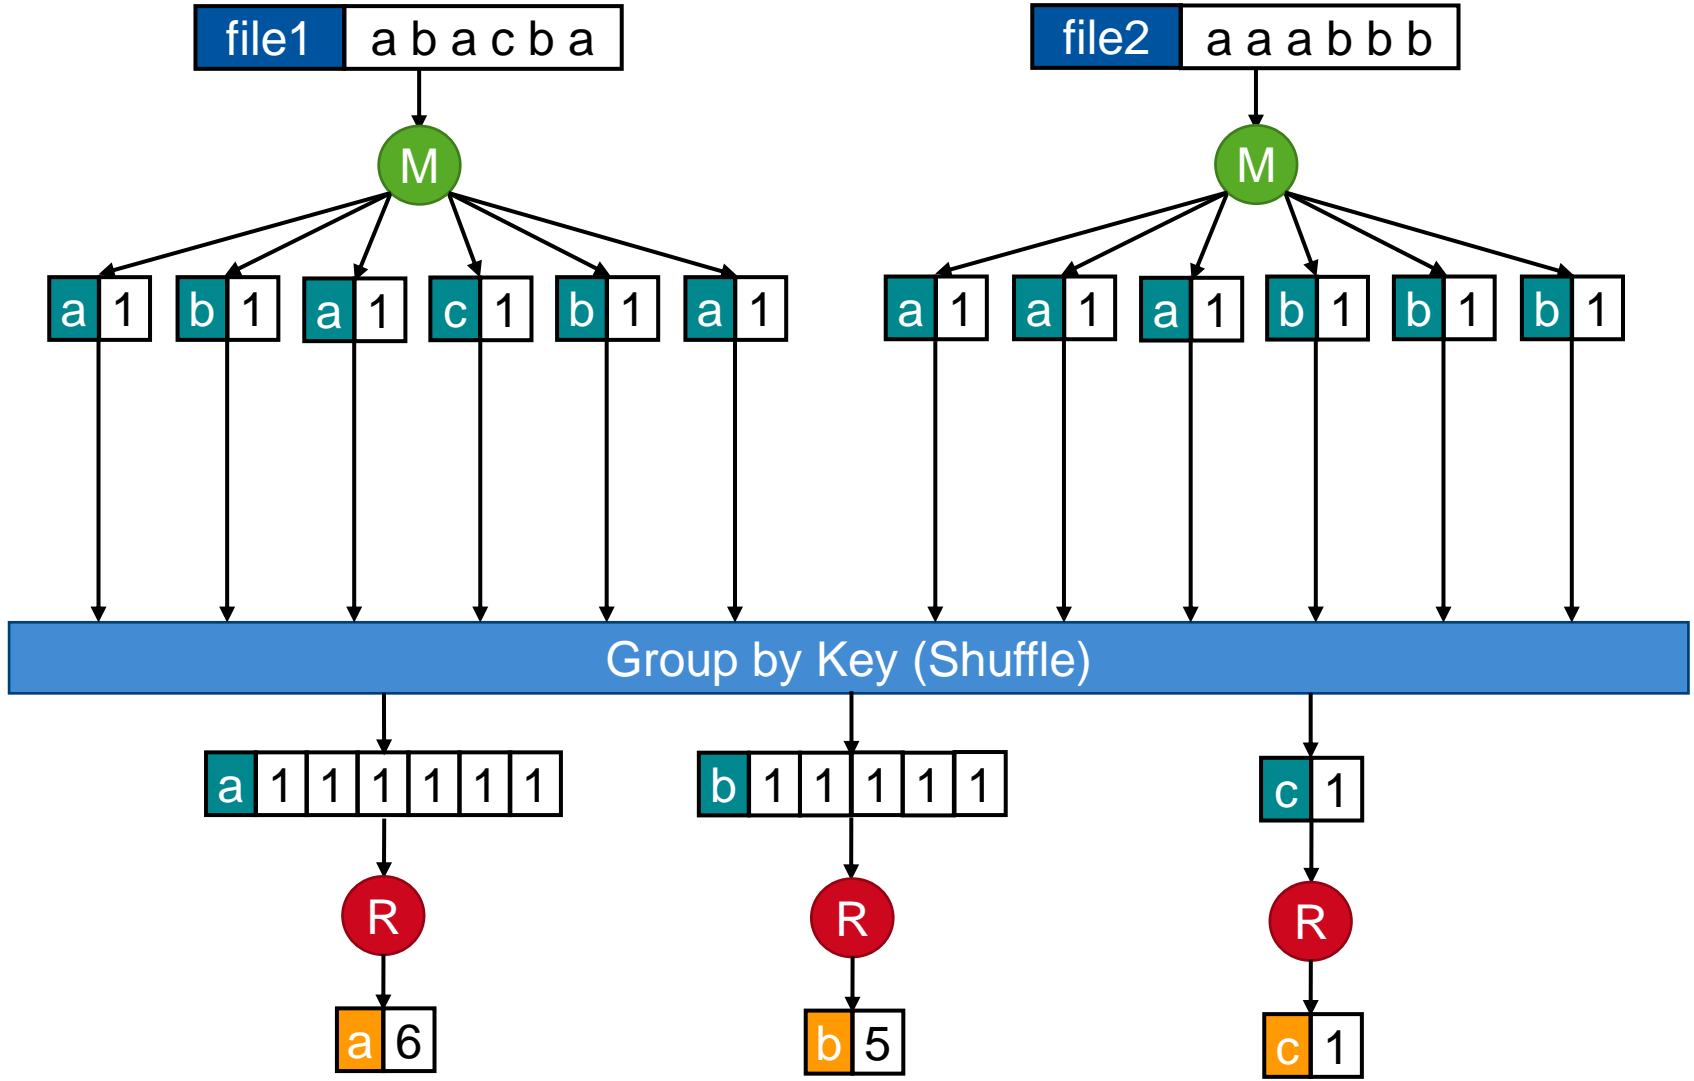


## Speculative Execution (2)

- Solution: Speculative execution
  - Detect tasks running at least one minute and with much less progress than the other tasks (on average) → Straggling task
  - If straggling task detected: Launch duplicate task as backup
  - Task is completed if primary or backup task completes, other instance is killed
- In practice: Significant reduction of total time
  - $\approx 30\%$  time reduction for a sorting program with 1,800 machines (Google)
- Drawback: Additional load on the cluster

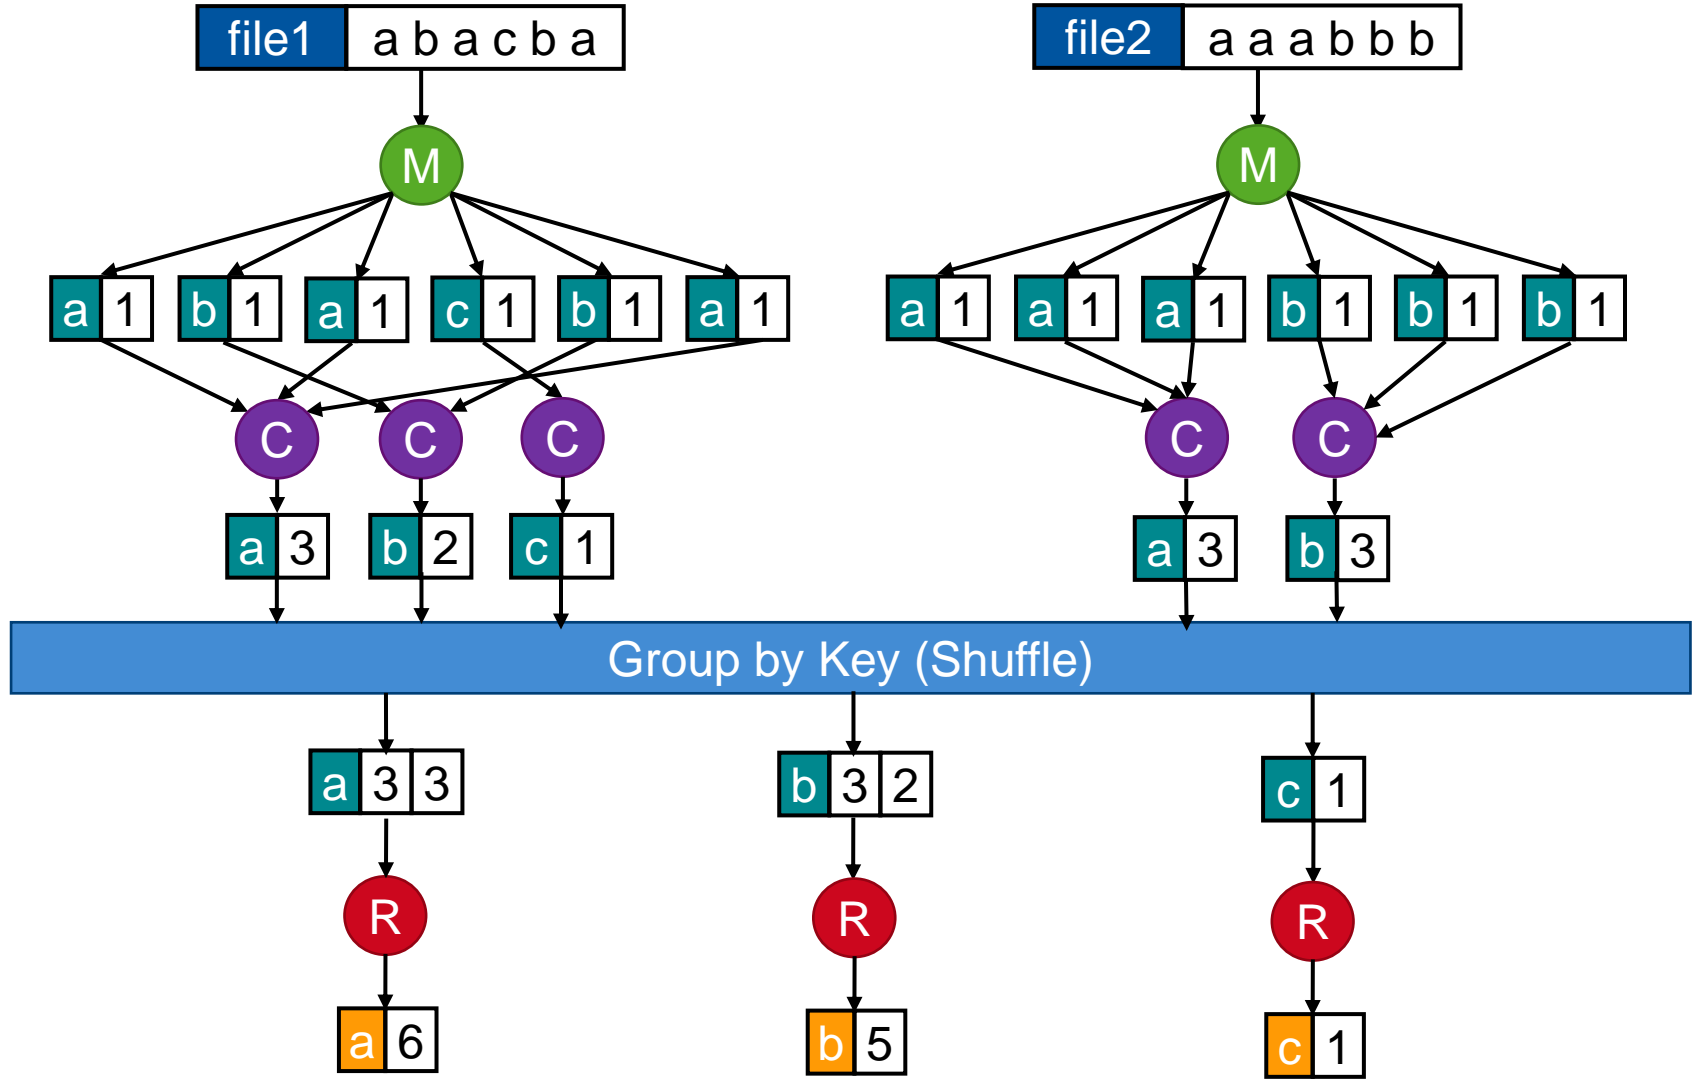


# Word Count Example – Revisited





# Combiner Function – Word Count



$$\text{Map}(k_1, v_1) \rightarrow \text{list}(k_2, v_2)$$

$$\text{Reduce}(k_2, \text{list}(v_2)) \rightarrow \text{list}(k_3, v_3)$$

## Combiner Function

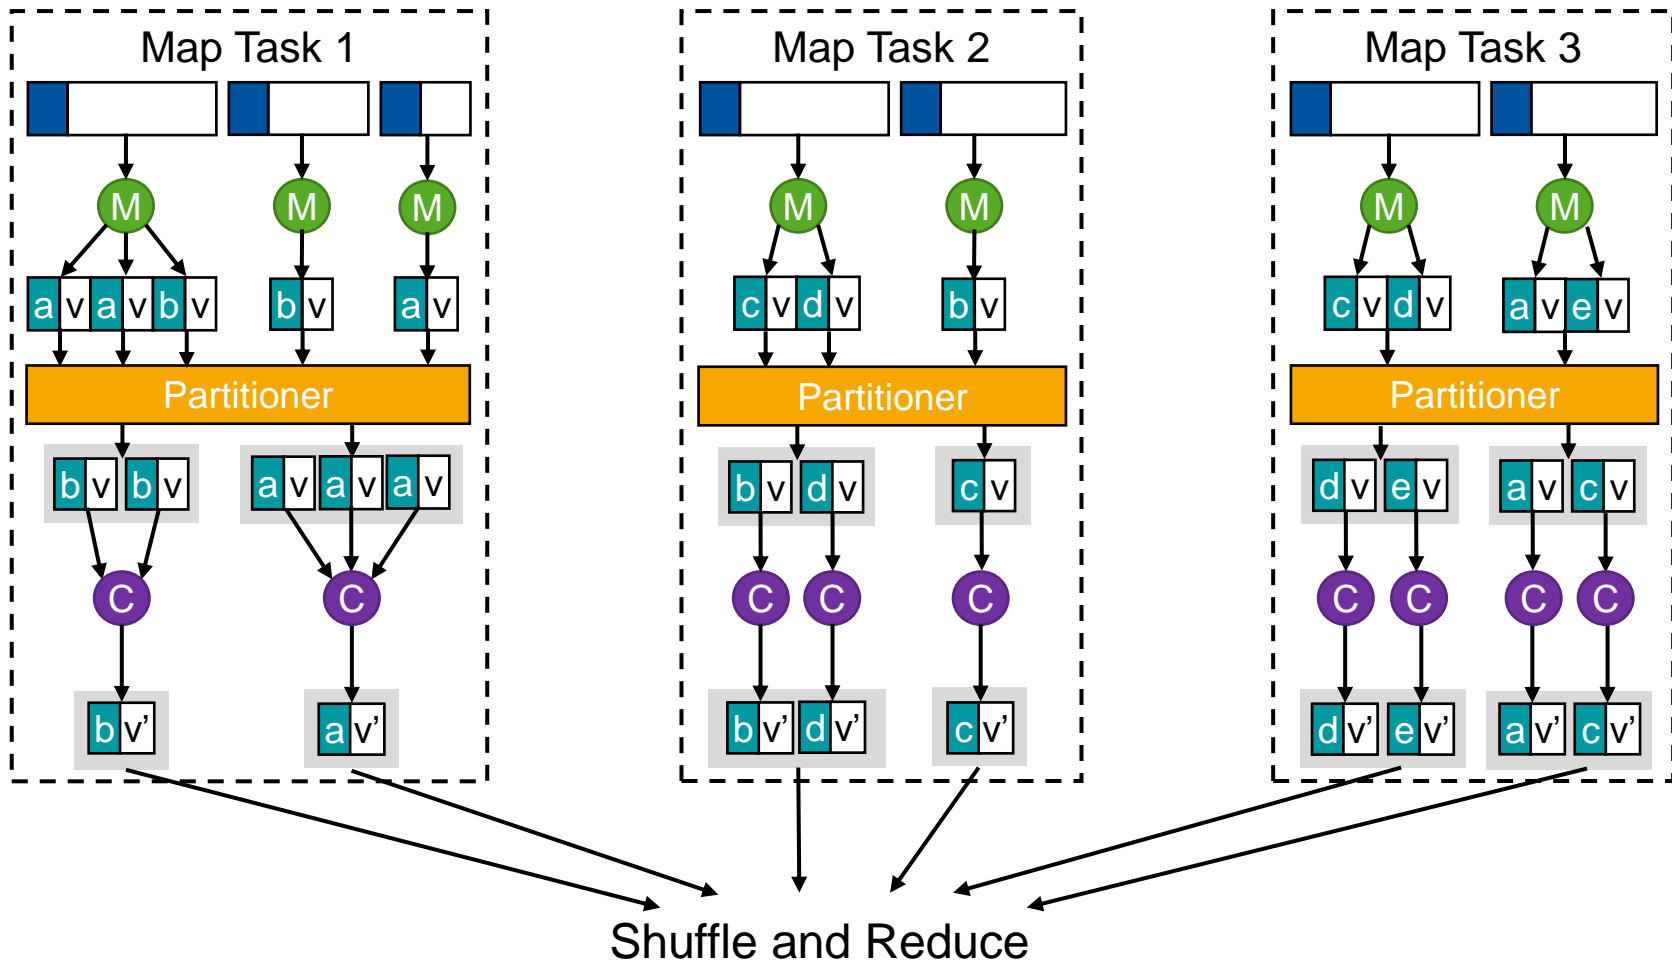
---

- Limit amount of data transferred between map and reduce task
- Idea: Aggregate (if possible) map output before reduction (“mini-reduce”)
- Optimization of Hadoop: Combiner function

$$\text{Combine}(k_2, \text{list}(v_2)) \rightarrow \text{list}(k_2, v_2)$$

- Runs on partitioned map output and aggregates (“combines”) list of items to a smaller list
- Runs directly after partitioning
- Potentially performed for each map task
- No guarantee how often it is called for a particular map task (zero or more invocations)
  - Computation should be independent of combiner invocations
  - The combiner is not allowed to change the key of the input KV pair, it can only aggregate the list of values.
- Word count example: Aggregate word counts for each map output

# MapReduce in Parallel with Combiner



# Job Submission with Combiner in Hadoop

---

## Job configuration and submission:

```
1  public class WordCount {
2      public static void main(String[] args) throws Exception {
3          Configuration conf = new Configuration();
4          Job job = Job.getInstance(conf, "word count");
5          job.setJarByClass(WordCount.class);
6
7          job.setMapperClass(WordCountMapper.class);
8          job.setCombinerClass(WordCountReducer.class);
9          job.setReducerClass(WordCountReducer.class);
10
11         job.setOutputKeyClass(Text.class);
12         job.setOutputValueClass(IntWritable.class);
13
14         FileInputFormat.addInputPath(job, new Path(args[0]));
15         FileOutputFormat.setOutputPath(job, new Path(args[1]));
16
17         System.exit(job.waitForCompletion(true) ? 0 : 1);
18     }
19 }
```