



Concepts and Models of Parallel and Data-centric Programming

MapReduce Design Patterns – Introduction

Lecture, Summer 2020


Simon Schwitanski
Dr. Christian Terboven

Outline

0. Organization
 1. Foundations
 2. Shared Memory
 3. GPU Programming
 4. Bulk-Synchronous Parallelism
 5. Message Passing
 6. Distributed Shared Memory
 7. Parallel Algorithms
 8. Parallel I/O
 - 9. MapReduce**
 10. Apache Spark
- a. MapReduce Programming Model
 - b. Parallelizing MapReduce
 - c. Hadoop Ecosystem
 - d. Hadoop Distributed File System
 - e. Yet Another Resource Negotiator
 - f. Comparison to Other Approaches
 - g. MapReduce Design Patterns**
 - a. Summarization Patterns
 - b. Filtering Patterns
 - c. Data Organization Patterns

Literature

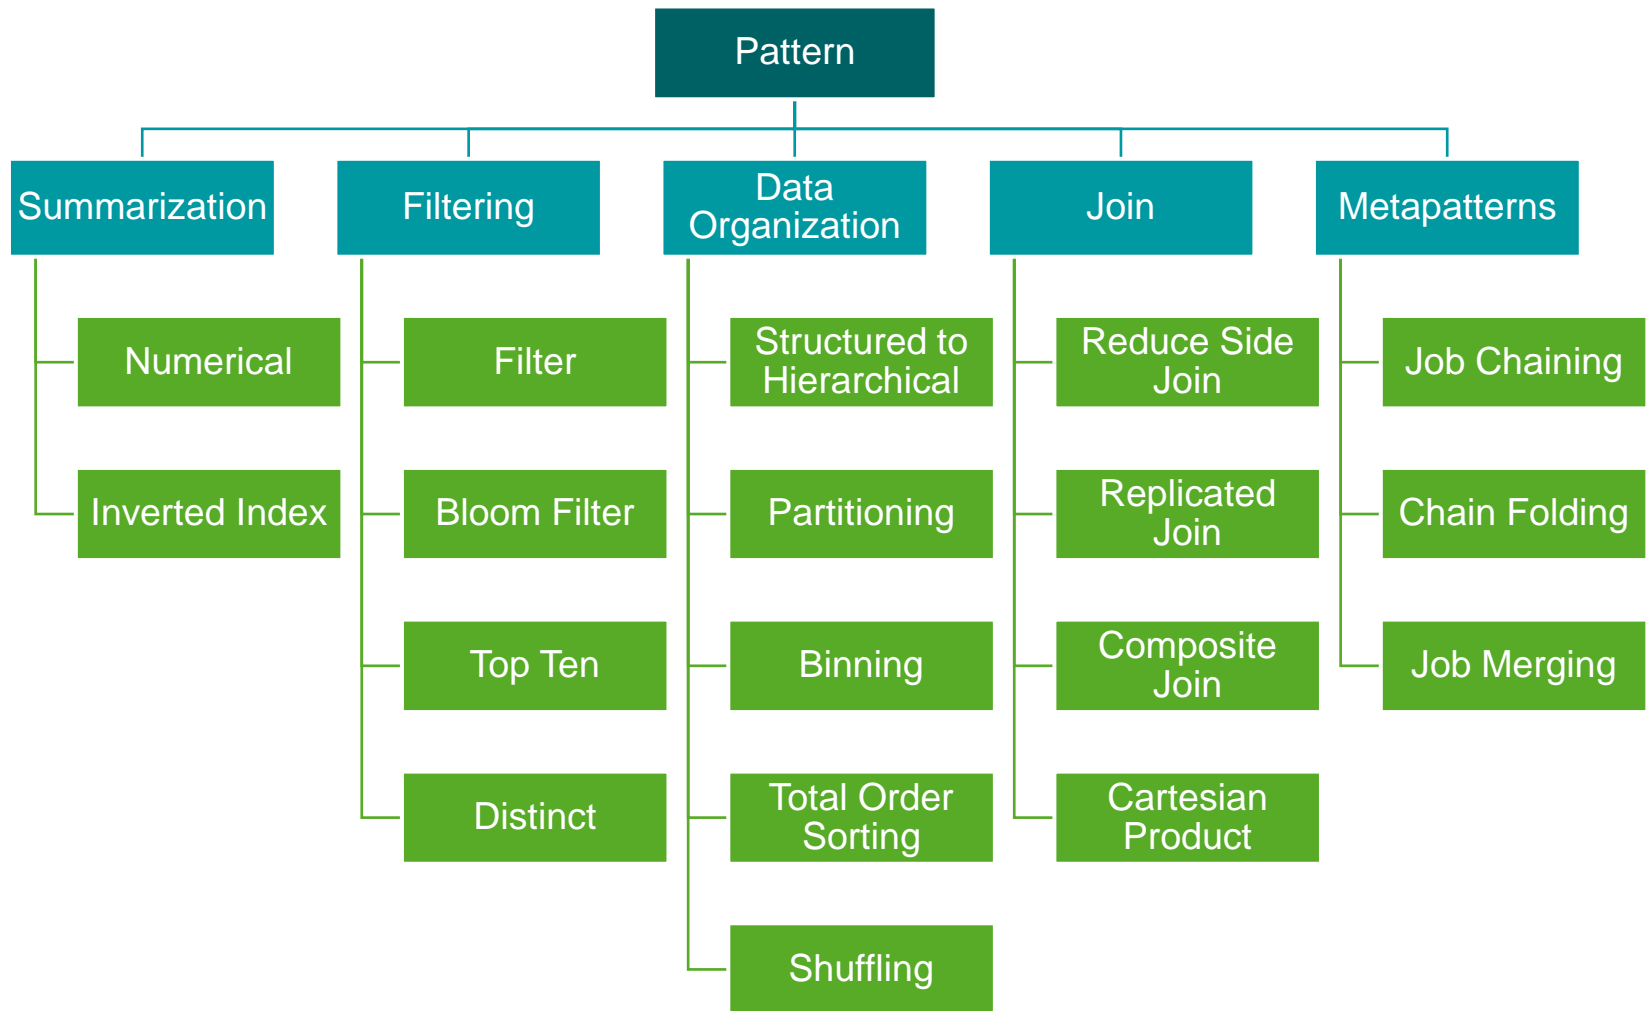
- Miner, Donald and Shook, Adam.
“MapReduce Design Patterns: Building Effective Algorithms and Analytics for Hadoop and Other Systems”, O'Reilly Media, 2012
- Slides based on this book
- Note: Reading this book is not required, neither for the lecture and exercises nor for the exam.



What is a Design Pattern?

- Blueprints for solving common problems in reusable and general way
- Popular SE design patterns in “Gang of Four” book
- Simplifies discussion with other developers and reading of other developer’s code
- MapReduce design patterns: Smaller problem space than SE patterns
- However: Design patterns incorporate knowledge of experienced developers
 - Helpful for new MapReduce programmers

Pattern Landscape



Pattern Landscape

