



Concepts and Models of Parallel and Data-centric Programming

MapReduce Design Patterns – Summarization Patterns

Lecture, Summer 2020

Simon Schwitanski
Dr. Christian Terboven

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. **MapReduce**
 - 10. Apache Spark
- a. MapReduce Programming Model
 - b. Parallelizing MapReduce
 - c. Hadoop Ecosystem
 - d. Hadoop Distributed File System
 - e. Yet Another Resource Negotiator
 - f. Comparison to Other Approaches
 - g. MapReduce Design Patterns
 - a. **Summarization Patterns**
 - b. Filtering Patterns
 - c. Data Organization Patterns

Numerical Summarization

- Calculation of aggregate statistical values
 - Counting, Minimum, Maximum, Average, Standard Deviation, ...
- Intent: Group records by *key* field, calculate aggregate value per group
- Motivation
 - Large datasets difficult to interpret and evaluate manually by a human
 - Numerical summarization gives abstracted overview of data
- Conditions for application
 - Data is numerical (or should be counted)
 - Data can be grouped by specific (key) fields

Numerical Summarization – Structure

- Mapper outputs key(s) to group by and values that should be aggregated (“summary field”)
- Reducer applies desired aggregation function to group values




Illustration: Miner, Donald and Shook, Adam. "MapReduce Design Patterns: Building Effective Algorithms and Analytics for Hadoop and Other Systems", p.16, O'Reilly Media, 2012

Numerical Summarization – Applications

- Word count
 - As previously discussed
- Record count
 - Count number of items in a certain interval (number of user comments per week etc.)
- Minimum, maximum and count of some particular event
- Average, median, standard deviation of numerical data
- Note: Numerical summarization similar to following SQL statement

SELECT MYFUNC(mycol) **FROM** mytable **GROUP BY** groupcol;

where MYFUNC is the desired aggregate function.

Numerical Summarization – Performance

- Perfectly suited for MapReduce (primarily designed for summarizations)
- Using a Combiner is crucial for large datasets
- Bottleneck can occur at Reducer
 - Imbalances in reduced data (much more intermediate data with a certain key than other keys) can lead to load imbalances

Working Data Set (1)

- Stack Exchange provides dump of posts, comments etc.
- Comments.xml

```
<row Id="6" PostId="13" Score="0" Text="Honestly, i was just about to ask this!" CreationDate="2014-02-11T23:46:29.183" UserId="8" />
```

- Posts.xml

```
<row Id="1" PostTypeId="1" AcceptedAnswerId="2" CreationDate="2014-02-11T22:24:09.530" Score="42" ViewCount="5664" Body="I'm making a simple Arduino web server [...]" OwnerUserId="3" Title="Is an Arduino capable of running 24/7?" AnswerCount="11" CommentCount="2" FavoriteCount="9" />
```

```
<row Id="2" PostTypeId="2" ParentId="1" CreationDate="2014-02-11T22:36:57.700" Score="57" Body="You shouldn't have any issues keeping it on all the time [...]" OwnerUserId="11" CommentCount="7" />
```

Working Data Set (2)

- Users.xml

```
<row Id="9" Reputation="131" CreationDate="2014-02-11T22:31:15.467"  
DisplayName="orangeocelot" LastAccessDate="2016-08-05T14:39:40.290"  
Location="Ireland" Views="4" UpVotes="7" DownVotes="0"  
AccountId="1391850" />
```

- Data is parsed in the examples using a helper function

```
public class MRDPUtils {  
    public static Map<String, String>  
        transformXmlToMap(String xml) {...}  
}
```

- First entry of map pair is the key, second entry the value

Numerical Summarization – Minimum Example (1)

- **Problem:** Given a list of user comments, determine for each user the earliest time he commented.
- Structure of comments:
`<row Id="6" PostId="13" Score="0" Text="Honestly, i was just about to ask this!" CreationDate="2014-02-11T23:46:29.183" UserId="8" />`
- One KV pair for each comment

Pseudocode:

```
map(String docid, String comment):  
    parsed = transformToMap(comment)  
    Emit(parsed.UserId, timestamp(parsed.CreationDate))
```

```
reduce(String userId, Iterator values):  
    long min = MAX_LONG  
    for each timestamp ts in values:  
        if ts < min:  
            min = ts  
    Emit(userId, min)
```

Numerical Summarization – Minimum Example (2)

- Structure of comments:

```
<row Id="6" PostId="13" Score="0" Text="Honestly, i was just about to ask  
this!" CreationDate="2014-02-11T23:46:29.183" UserId="8" />
```

Mapper:

```
1 public static class EarliestCommentMapper extends  
2     Mapper<Object, Text, Text, LongWritable> {  
3     // Output key and value Writable  
4     private Text outUserId = new Text();  
5     private LongWritable outDateLong = new LongWritable();  
6  
7     // This object will format the creation date string into a Date object  
8     private final static SimpleDateFormat frmt = new SimpleDateFormat(  
9         "yyyy-MM-dd'T'HH:mm:ss.SSS");
```

Numerical Summarization – Minimum Example (3)

Mapper (continued):

```
10  @Override
11  public void map(Object key, Text value, Context context)
12      throws IOException, InterruptedException {
13      // Parse the input string into a nice map
14      Map<String, String> parsed =
15          MRDPUtils.transformXmlToMap(value.toString());
16
17      String userId = parsed.get("UserId");    // extract user id
18      String strDate = parsed.get("CreationDate"); // extract creation date
19      Date creationDate = frmt.parse(strDate); // translate to Date object
20
21      // Write out the user ID with creation date (as long)
22      outUserId.set(userId);
23      outDateLong.set(creationDate.getTime());
24      context.write(outUserId, outDateLong);
25  }
26 }
```

Numerical Summarization – Minimum Example (4)

Reducer:

```
1  public static class EarliestCommentReducer extends
2      Reducer<Text, LongWritable, Text, LongWritable> {
3      private LongWritable result = new LongWritable();
4
5      @Override
6      public void reduce(Text key, Iterable<LongWritable> values,
7                          Context context)
8          throws IOException, InterruptedException {
9          // Set initial value
10         long min = Long.MAX_VALUE;
11
12         // Iterate through all input values for this key
13         for (LongWritable val : values) {
14             long longValue = val.get();
15             if (longValue < min)
16                 min = longValue;
17         }
18
19         // Write out result
20         result.set(min);
21         context.write(key, result);
22     }}
```

Numerical Summarization – Minimum Example (5)

Job configuration:

```
1  public static void main(String[] args) throws Exception {
2      Configuration conf = new Configuration();
3
4      Job job = Job.getInstance(conf, "StackOverflow Comment Date Min");
5      job.setJarByClass(EarliestCommentDate.class);
6
7      job.setMapperClass(EarliestCommentMapper.class);
8      job.setReducerClass(EarliestCommentReducer.class);
9      job.setOutputKeyClass(Text.class);
10     job.setOutputValueClass(LongWritable.class);
11     FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
12     FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
13     System.exit(job.waitForCompletion(true) ? 0 : 1);
14 }
```

Numerical Summarization – Running on Cluster

- “Our” Hadoop cluster 2018: 11 CLAIX-2016-MPI nodes
 - 1 master, 10 workers
 - 2 x Intel Xeon E5-2650v4, 12 cores each
 - 128 GB RAM per node
 - Interconnect: Intel Omni-Path, up to 100 Gbit/s network bandwidth
 - Access to files via Lustre file system, no local storage → 2020: Local storage
- You will submit and run example tasks on our custom Hadoop cluster in the exercises.
- Dataset: All comments on StackOverflow (XML file)
 - Entries: 66,432,644
 - Size: 17 GB

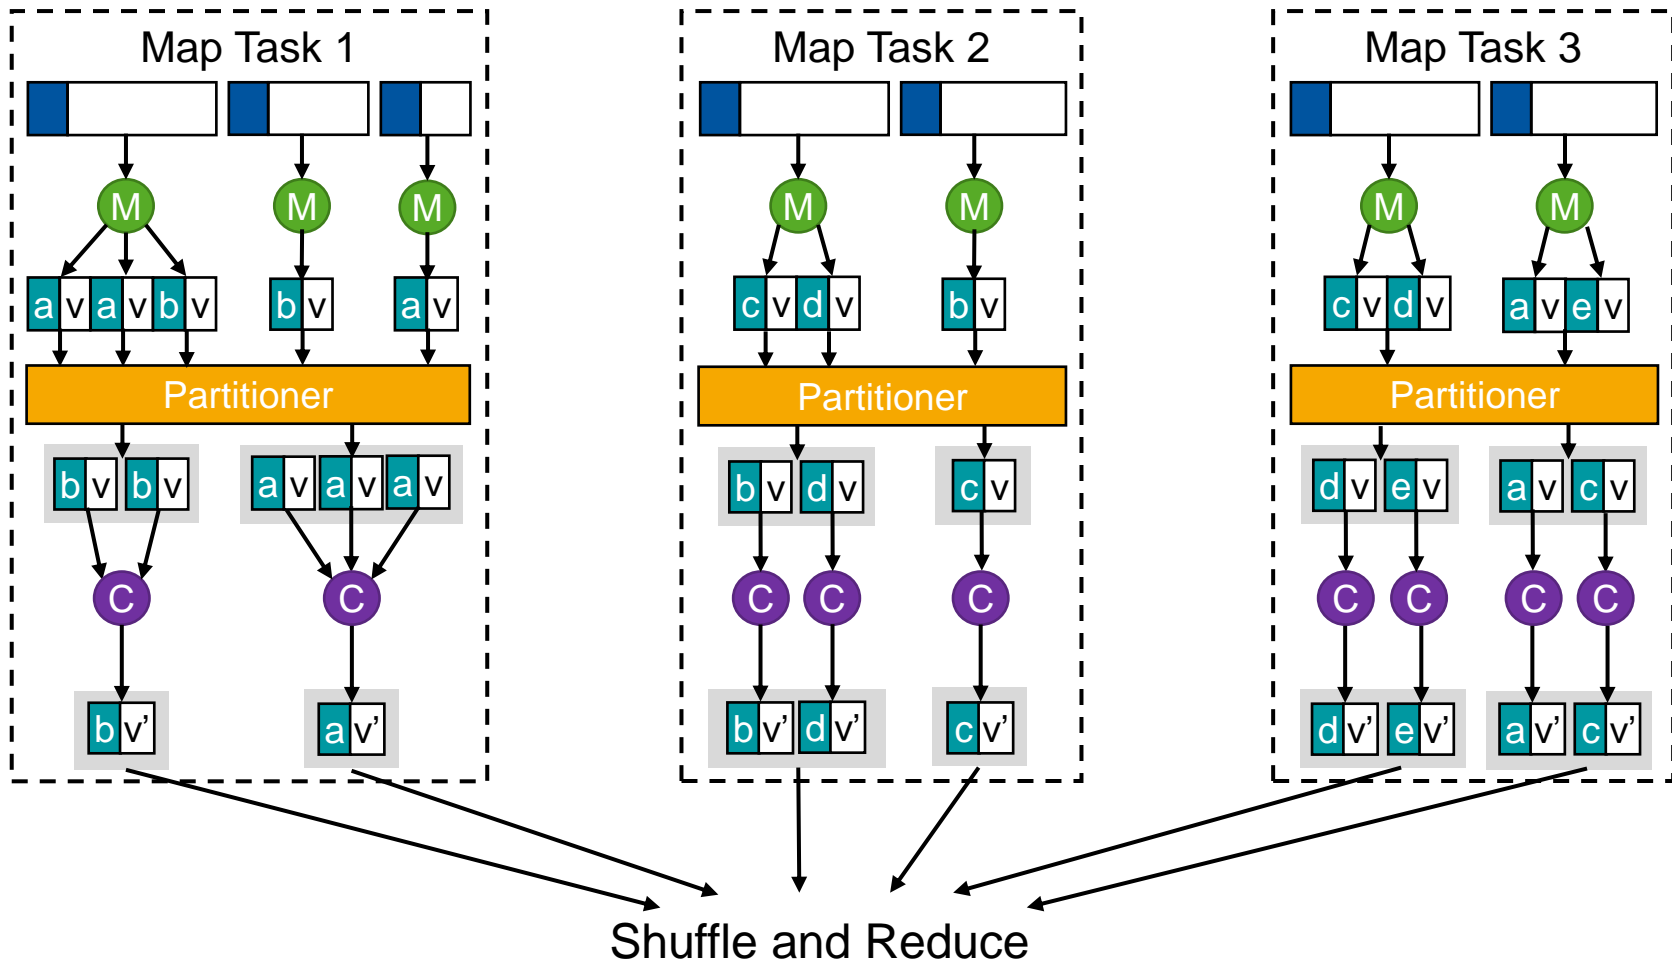
Numerical Summarization – Running on Cluster

Output Data:	User ID	Earliest Comment (Timestamp)
	10	1392258686390
	10003	1433121838880
	10007	1432600989783
	10017	1431534771737

Output Log:

- Map input records=66432644
- Map output records=65722799
- Map output bytes=1019807085
- Map output materialized bytes=1151291983
- Input split bytes=15458
- Combine input records=0
- Combine output records=0
- Reduce input groups=2502332
- Reduce shuffle bytes=1151291983 \approx 1 GB
- Reduce input records=65722799
- Reduce output records=2502332

Combiner Function – Reminder



$$\text{Combine}(k_2, \text{list}(v_2)) \rightarrow \text{list}(k_2, v_2)$$

Numerical Summarization – Minimum with Combiner

- Job configuration:

```
1  public static void main(String[] args) throws Exception {
2      Configuration conf = new Configuration();
3
4      Job job = Job.getInstance(conf, "StackOverflow Comment Date Min");
5      job.setJarByClass(EarliestCommentDate.class);
6
7      job.setMapperClass(EarliestCommentMapper.class);
8      job.setCombinerClass(EarliestCommentReducer.class);
9      job.setReducerClass(EarliestCommentReducer.class);
10     job.setOutputKeyClass(Text.class);
11     job.setOutputValueClass(LongWritable.class);
12     FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
13     FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
14     System.exit(job.waitForCompletion(true) ? 0 : 1);
15 }
```

Numerical Summarization – Running on Cluster

Output Log without Combiner:

Map input records=66432644
Map output records=65722799
Map output bytes=1019807085
Map output materialized bytes=1151291983
Input split bytes=15458
Combine input records=0
Combine output records=0
Reduce input groups=2502332
Reduce shuffle bytes=1151291983 \approx 1 GB
Reduce input records=65722799
Reduce output records=2502332

Output Log with Combiner:

Map input records=66432644
Map output records=65722799
Map output bytes=1019807085
Map output materialized bytes=204036720
Input split bytes=15458
Combine input records=65722799
Combine output records=11569585
Reduce input groups=2502332
Reduce shuffle bytes=204036720 \approx 200 MB
Reduce input records=11569585
Reduce output records=2502332

- Network I/O reduced by factor 5 due to combiner
- Effort: Adding one LOC (set combiner in job configuration)
- However: Nearly no impact on runtime
 - Reason: Network interconnect extremely fast, network I/O makes no difference for 200 MB / 1 GB, data set too small

Numerical Summarization – Average (1)

- Calculation of minimum, maximum and count simple
- Can we do the same with a measure like average?
- First answer: Yes!
- **Example:** Get average comment length posted per hour of day.

```
map(String docid, String comment):  
    parsed = transformToMap(comment)  
    Emit(getHour(parsed.CreationDate), length(parsed.Text))
```

```
reduce(String hour, Iterator values):  
    int count = 0  
    int sum = 0  
    for each value v in values  
        count++  
        sum += v  
    Emit(hour, ((float) sum) / count)
```

- Can we use this Reducer as Combiner for average calculation?

Numerical Summarization – Average (2)

- No, because the average operation is *not associative*.
- Average of two values: $avg(a, b) = \frac{a+b}{2}$
- Associativity: $avg(avg(a, b), c) = avg(a, avg(b, c))$
- Counterexample: $a = 2, b = 6, c = 10$
 - $avg(avg(2, 6), 10) = avg(4, 10) = 7$
 - $avg(2, avg(6, 10)) = avg(2, 8) = 5$
- Thus: Associativity **not** fulfilled for average operation.
- Using the previously defined Reducer as Combiner will not work.

Numerical Summarization – Average (3)

- Average operation for a set D of values: $avg(D) = \frac{sum(D)}{count(D)}$
- Idea: Mapper and Reducer output *count* sent along with *average* of values → Reproducing *sum* for recomputing *average* possible:

$$sum(D) = avg(D) \cdot count(D)$$

```
map(String docid, String comment):  
    parsed = transformToMap(comment)  
    Emit(getHour(parsed.CreationDate), new Pair(1, length(parsed.Text)))
```

```
reduce(String hour, Iterator values):  
    float sum = 0;  
    float count = 0;  
    for each pair (countold, average) in values:  
        sum += countold * average  
        count += countold  
    Emit(hour, new Pair(count, sum / count))
```

← Reproduce *sum* using “previous”
count and *average*

- Now, Reducer can be also used as Combiner.

Numerical Summarization – Average (4)

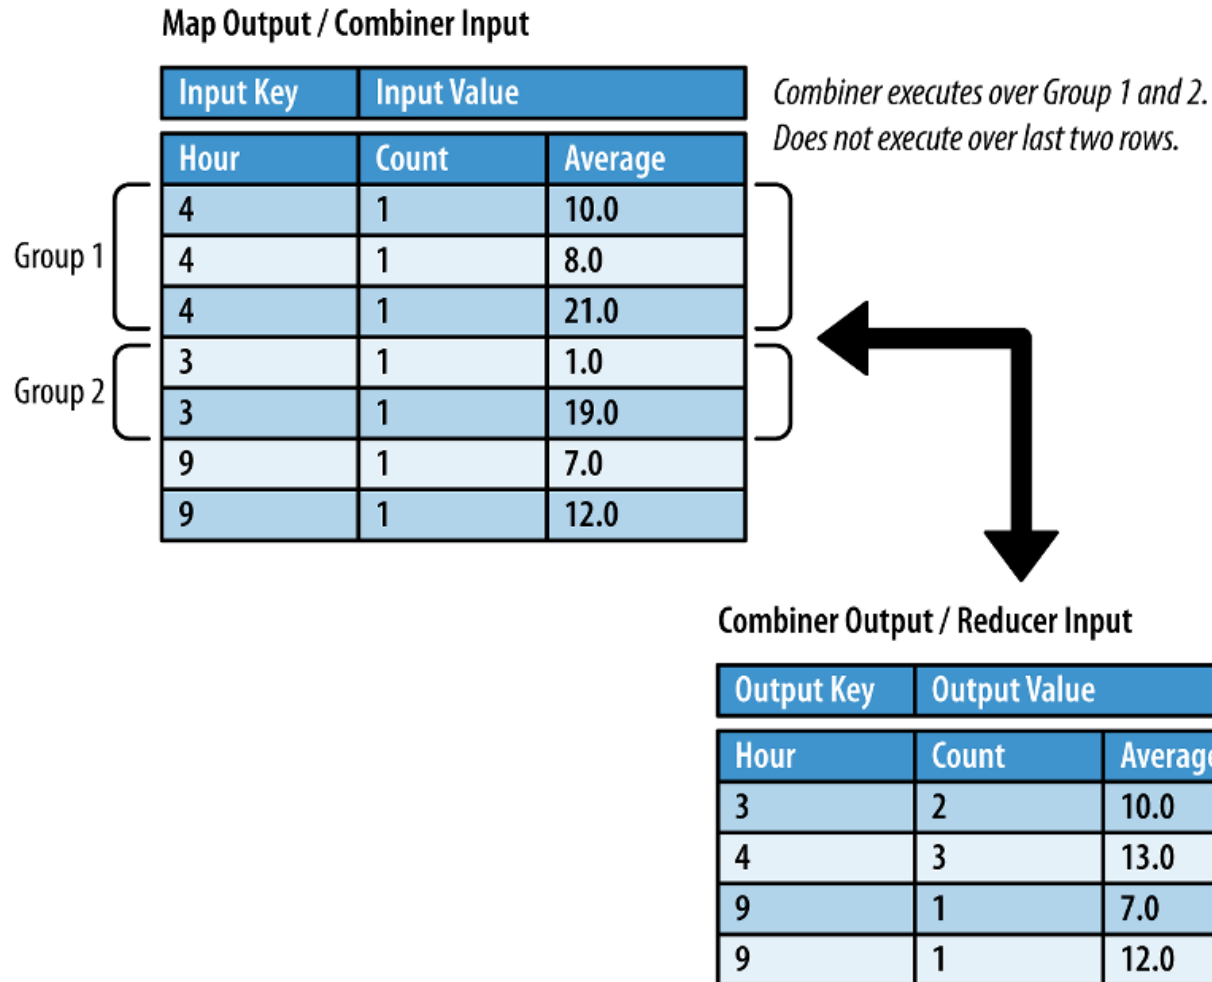


Illustration: Miner, Donald and Shook, Adam. "MapReduce Design Patterns: Building Effective Algorithms and Analytics for Hadoop and Other Systems", p.22, O'Reilly Media, 2012