



Concepts and Models of Parallel and Data-centric Programming

MapReduce Design Patterns – Data Organization Patterns

Lecture, Summer 2020

Simon Schwitanski
Dr. Christian Terboven

Outline

- 0. Organization
 - 1. Foundations
 - 2. Shared Memory
 - 3. GPU Programming
 - 4. Bulk-Synchronous Parallelism
 - 5. Message Passing
 - 6. Distributed Shared Memory
 - 7. Parallel Algorithms
 - 8. Parallel I/O
 - 9. **MapReduce**
 - 10. Apache Spark
- a. MapReduce Programming Model
 - b. Parallelizing MapReduce
 - c. Hadoop Ecosystem
 - d. Hadoop Distributed File System
 - e. Yet Another Resource Negotiator
 - f. Comparison to Other Approaches
 - g. MapReduce Design Patterns
 - a. Summarization Patterns
 - b. Filtering Patterns
 - c. Data Organization Patterns**

Data Organization Patterns (1)

- Reorganization of data often crucial for performance
- Partitioning, sharding and sorting can significantly affect runtime
- Structured to Hierarchical pattern
 - Create new records from data which has originally another structure
 - Migration from RDBMS to MapReduce: Transforming row-based data into hierarchical format like JSON or XML → More flexible, avoids joins
 - Example: Grouping a StackOverflow question together with answers and all comments

Data Organization Patterns (2)

- Partitioning pattern
 - Move records into categories (“partitioning”)
 - Uses partitioner of MapReduce framework and identity mapper and reducer
 - Write custom partitioner to change partitioning
 - Example: Partition users by last access date

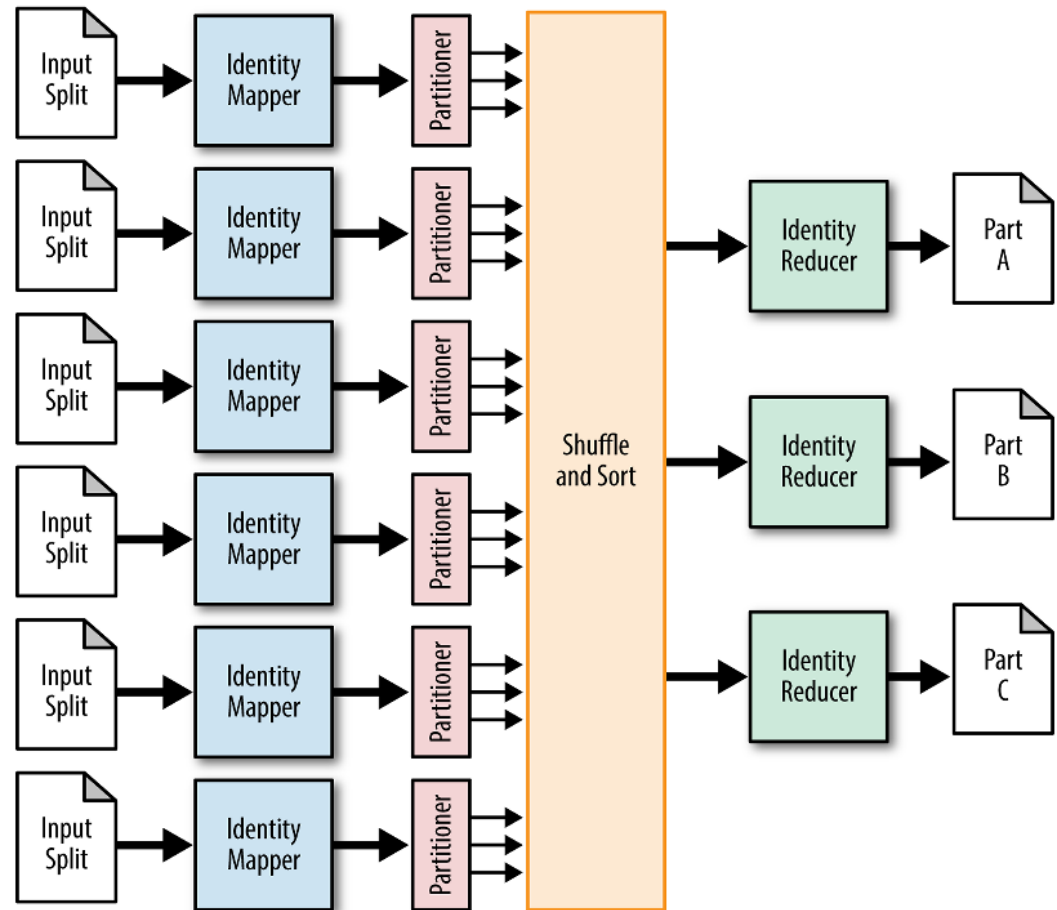


Illustration: Miner, Donald and Shook, Adam. “MapReduce Design Patterns: Building Effective Algorithms and Analytics for Hadoop and Other Systems”, p.84, O'Reilly Media, 2012

Data Organization Patterns (3)

- Binning pattern
 - Similar to partitioner pattern, but binning takes place directly in the mapper
 - No reduce phase
 - Each mapper produces one output file for each bin (huge number of files depending on number of bins and map tasks)

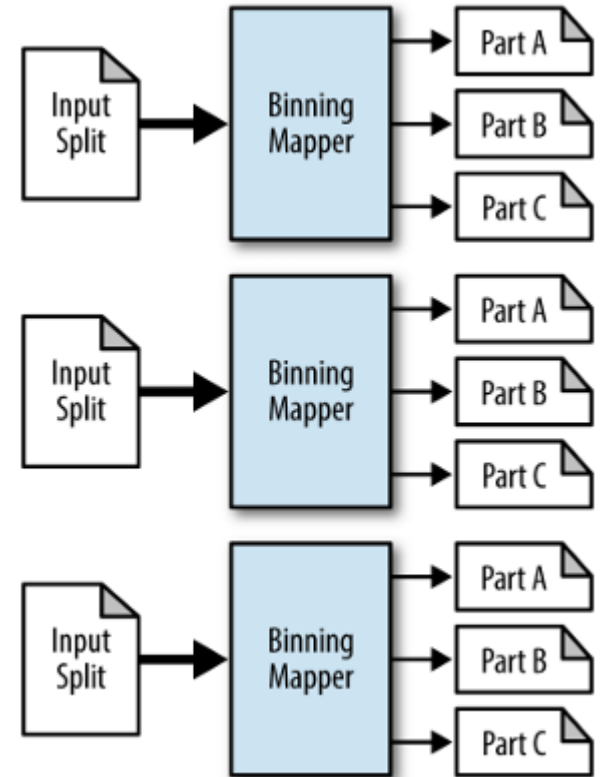


Illustration: Miner, Donald and Shook, Adam.
"MapReduce Design Patterns: Building Effective Algorithms and Analytics for Hadoop and Other Systems", p.89, O'Reilly Media, 2012

Data Organization Patterns (4)

- Total Order Sorting pattern
 - Sorting data in parallel on a given sort key
 - Sorting in MapReduce not easy, because distributing total order sorting between tasks is difficult
 - Each Reducer sorts data by key (data sorted per partition), but no global sorting across all data
 - Main idea: Write custom partitioner that partitions data by sort key: Lowest range to first reducer, next range to second reducer, ...
 - Challenge: Find suitable ranges for values
- Shuffling pattern
 - Randomizing set of records (opposite of total order sorting pattern)
 - Idea: Mapper outputs record with a random key, reducer sorts random keys and thus further randomizes data
 - Application: Anonymization

Concluding Remarks

Limitations of Hadoop MapReduce

- No support of caching data in memory, only disk-based
 - Each MapReduce job: Read from local disk, processing, write back to the file system
 - MapReduce job working on previous output has to read from local disk again
 - Serious performance penalty for iterative algorithms
 - Problem: Most data mining and machine learning algorithms are iterative
- Focus on batch processing, interactive analysis inefficient
- Inflexibility: Algorithms have to be broken down to *Map* and *Reduce* functions

Alternatives

- Specialized frameworks
 - HaLoop [1]: Iterative MapReduce interface (loop-aware scheduler and caching)
 - Pregel [2]: Model for iterative graph applications (keeps intermediate data in memory)
- Generalized framework: Apache Spark [3]
 - Implements a caching feature → Allows to reuse the outputs of previous operations
 - Beneficial for iterative algorithms
 - Interactive analysis supported → Continuously querying data possible
 - Processing of real-time data → Data from real time event streams (e.g., Twitter)

[1] Yingyi Bu, Bill Howe, Magdalena Balazinska, Michael D. Ernst. "HaLoop: Efficient Iterative Data Processing on Large Clusters" PVLDB 3(1): 285-296 (2010)

[2] Grzegorz Malewicz, Matthew H. Austern, Aart J. C. Bik, James C. Dehnert, Ilan Horn, Naty Leiser, Grzegorz Czajkowski. "Pregel: a system for large-scale graph processing" SIGMOD Conference 2010: 135-146

[3] Matei Zaharia, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, Ion Stoica. "Spark: Cluster Computing with Working Sets" HotCloud 2010

What you have learnt

What you have learnt

- MapReduce design patterns: Blueprints for solving common problems in reusable and general way
 - Summarization
 - Filtering
 - Data Organization
- Limitations of MapReduce: Only support for batch processing
 - Iterative jobs often limited by disk I/O, no caching supported
 - No interactive analyses possible
- Generalization of MapReduce approach: Apache Spark